

FuzzPilot: Extending Off-Hot-Path LLM Control to Structured Text Parsers via Data-as-Recipe and Micro-Campaign Validation

Zhiyi Yao

Qingdao University of Technology
202202030221@stu.qut.edu.cn

May 24, 2026

Abstract

Recent work has shown that placing a large language model *off the hot path* of a greybox fuzzer can preserve mutation throughput while still letting the model contribute strategy: G²FUZZ, in particular, has the LLM synthesize Python input generators that are then mutated by AFL++ [14]. That paradigm targets *non-textual binary formats* such as JPEG and TIFF, where input generators in code form are a natural intermediate representation. We adapt the off-hot-path paradigm to *structured text parsers* by building three reusable pieces of infrastructure on top of AFL++: (i) a **recipe-guided custom mutator** that dispatches schema-validated, cacheable mutation strategies as native code on the hot path, with no LLM call per mutation; (ii) a **micro-campaign validation gate** that scores each candidate recipe in an isolated N -second AFL++ run against an explicit reward function over new edges, paths and crashes, and refuses to promote candidates with non-positive reward; (iii) a **static-analysis context channel** (Ghidra headless) that feeds both the agent blackboard and the in-loop AFL dictionary from a single cached extraction. An LLM is hooked into this pipeline as a plateau-triggered proposal source, but the preprint’s claims are anchored to the mutator-and-gate framework rather than to LLM-derived recipes. This preliminary report evaluates the system on a single structured-text target, cJSON, with $N=5$ repetitions per mode for the two main configurations (**baseline-afl** and **full-agent**, each at 14,400s wall-clock); we do *not* report results on additional targets here, and we *do not claim* generalization to other parsers (JSON, XML, SQL beyond cJSON, libpng, etc.). Within that scope, three findings and one negative result stand out. **First**, end-to-end throughput parity holds: **full-agent** median `execs_per_sec` is $\approx 1.06\times$ baseline at $N=5$ (§6.3). **Second**, **full-agent** shows a descriptive 45.3% shorter median plateau than **baseline-afl** (1,384s vs 2,532s; §6.2); both modes reach the same 269-edge ceiling that cJSON’s harness exposes. At $N=5$ this delta is *not* statistically significant (Mann–Whitney $U=8$, $p=0.42$; Vargha–Delaney $A_{12}=0.68$) and we report it as a point estimate rather than a significance claim. **Third**, the gate did not promote any LLM-derived recipe on cJSON: across the 5 full-agent runs the micro-campaign gate evaluated 20 LLM candidates and *promoted zero*, because cJSON’s edge space has already been exhausted by the parent corpus at plateau time, leaving the gate’s reward at $R=0$ on every candidate (§6.4). The plateau-duration delta is therefore attributable to the recipe-guided mutator running the controller’s *default DictionaryAgent* recipe rather than to LLM-promoted recipes; whether this attribution survives a non-saturated-target evaluation is the central open question for the venue version. **Negative result**: the LLM proposal layer’s contribution cannot be empirically demonstrated on a saturated target; what we *can* demonstrate is that the infrastructure path behaves predictably ($43/45 = 95.6\%$ schema-valid responses, zero fallbacks, total inference cost \$0.002 per 4-hour campaign). The full code, configurations, decision logs, and raw artifacts are released to support replication and to enable a non-saturated-target evaluation in the venue paper.

1 Introduction

Greybox fuzzers such as AFL++ achieve their effectiveness by running target programs hundreds of thousands of times per second; on modern hardware the inner mutation loop sustains 10^5 to 10^6 executions per second on small parsers [8]. Recent work integrates large language models (LLMs) into this loop to inject semantic awareness that bit-level mutators cannot provide. The dominant design in this space — exemplified by LLAMAFUZZ [21], FuzzCoder [15], and Fuzz4All [24] — invokes the model *within* the mutation loop, paying the model’s per-call latency on every mutation. A recent semantic-feedback design pays a slightly lower price by running an LLM-driven helper fuzzer in parallel at roughly 250 queries per hour [12], but the LLM remains a steady-state cost during fuzzing.

A second, less explored line of work removes the model from the mutation loop entirely. G²FUZZ [14] invokes the LLM only when the underlying fuzzer fails to discover new coverage; the LLM emits a Python script (an “input generator”) which is then executed to produce many candidate inputs that AFL++ mutates as usual. We refer to this as the *off-hot-path* paradigm: the LLM contributes to strategy at coarse intervals and never participates in per-mutation work. G²FUZZ demonstrates that off-hot-path placement is achievable in practice on *non-textual* input domains such as JPEG, TIFF, and MP4, where Python scripts are a natural way to instantiate format-specific constructors.

Off-hot-path placement raises three open design questions that the non-textual-binary setting in prior work does not surface. (Q1) When the target consumes structured *text* — JSON objects, XML documents, SQL statements — is a generator script still the right intermediate representation, given that the LLM could instead emit a *mutation strategy* that the existing mutator dispatches without running new code at all? (Q2) When the LLM proposes a strategy, how do we know it actually helps before committing it to a multi-hour campaign? (Q3) Where does the structured prompt context come from when the target is a stripped binary or its source is unavailable?

This paper presents FuzzPilot, which answers (Q1)–(Q3) for the structured-text-parser domain.

Contributions. FuzzPilot’s contributions are deliberately incremental over the off-hot-path paradigm: we adopt it as the contract and refine three aspects that the prior work either does not explore or solves only for the non-textual case.

- C1 A data-as-recipe** intermediate representation that replaces LLM-generated code with a schema-validated description of mutation strategy — a weighted distribution over seven mutation operators, plus focus/protect byte ranges, dictionary tokens, and a corpus selector — dispatched by a native AFL++ custom mutator. This side-steps the runtime-safety concern of executing model-emitted code, makes proposals cacheable and version-able, and lets the mutator hot path remain pure native code at AFL++ speed.
- C2 A micro-campaign validation gate** that scores each LLM proposal in an N -second isolated AFL++ run over a corpus snapshot and promotes only the top- K by a reward function over new edges and paths. The gate gives the controller empirical evidence about each proposal before it enters the main loop, and bounds the worst-case cost of a poor proposal.
- C3 A static-analysis context channel** built on Ghidra headless. Once per binary, a custom post-script (`FuzzPilotGhidraExtract.java`) extracts six categories of signal: a danger-scored ranking of risky functions, magic tokens (strings and 4-byte ASCII constants), comparison constants from operand scalars, branch constraints recovered by lookahead before conditional jumps, decompiled C for the top-ranked functions, and struct definitions. The signal

feeds two consumers simultaneously — the agent blackboard (for LLM prompts) and the micro-campaign’s AFL dictionary (for direct mutation guidance) — with results cached as `base_intelligence.json` so subsequent runs on the same binary do not re-analyze. The ablation E2c (§6.5) jointly quantifies the loss of both consumers; an isolated binary-target evaluation is deferred to the venue paper.

C4 An **end-to-end auditable decision log** (blackboard → proposal → recipe → micro-campaign reward → promotion outcome) and a Docker-based reproducibility artifact, addressing the long-standing reviewer concern that LLM-driven fuzzing results are hard to re-execute.

Scope. This is an arXiv preliminary report. We evaluate on a *single* target, cJSON, at $N=5$ repetitions per mode for the two main configurations (**baseline-afl** and **full-agent**). The microbench (E3) measures off-loop mutator dispatch and is target-independent. We make *no* empirical claim about libpng, libxml2, sqlite3, openssl, or any non-textual binary format here; multi-target evaluation, formal significance testing at higher repetition counts, and a head-to-head against G²FUZZ on its non-textual benchmark are deferred to the full venue version (Paper 3). The **rule-only** ablation E2a and **no-mutator** ablation E2b (3 preliminary runs each) are summarized in §6.5; the **no-static-analysis** (E2c) arm completed retry on the experiment host; results are reported in §6.5 (median plateau 234s, bimodal at $N=3$). The LLM proposal layer hooks into the architecture, but the gate produces a null result on cJSON (§6.4); a non-saturated-target evaluation that exercises the LLM layer is the most important single piece of future work. §8 states this scope honestly.

Roadmap. §2 surveys the design space and positions FuzzPilot relative to G²FUZZ and the in-loop line of work. §3 presents the architecture. §4 details the recipe representation and contrasts data-as-recipe with code-as-generator. §5 specifies the plateau detector, the Ghidra context channel, and the micro-campaign protocol. §6 reports preliminary results on cJSON. The per-run decision trail (events log, recipe proposals, micro-campaign reward, promotion outcome) is recoverable from `results/.../runs/p1_e1_cjson_full-agent.r*/events.jsonl`. §7 relates FuzzPilot to the broader LLM-fuzzing literature. §8 states limitations and threats to validity.

2 Background and the Off-Hot-Path Paradigm

2.1 The AFL++ mutation loop and its throughput budget

AFL++ implements a deterministic-then-havoc mutation loop [8]: each cycle takes an input from the queue, applies a sequence of bit/byte mutations, executes the target via `fork+execve` or a persistent harness, and updates the coverage bitmap. On small parsers such as cJSON our 4-core x86_64 host sustains roughly 13,000 executions per second per worker. Any redesign that interposes a non-trivial computation on this loop caps end-to-end throughput at the rate of that computation: inserting a single 100 ms LLM call per mutation drops throughput by three orders of magnitude. Existing LLM-in-the-loop designs accept this cost in exchange for semantic input quality.

2.2 In-loop LLM fuzzers

LLAMAFUZZ [21] fine-tunes an LLM on seed-mutation pairs and queries it as part of the mutation step; the model output is then handed to AFL++. FuzzCoder [15] casts mutation as a sequence-to-sequence task and predicts mutation positions and strategies per input. Fuzz4All [24] drives an

LLM as the primary input generator and mutation engine across multiple input languages. The hybrid design of Lin et al. [12] reduces the per-mutation cost by isolating LLM work in a helper fuzzer at roughly 250 queries per hour, with a syntax-validator-and-repair filter, while a master fuzzer runs AFL++ at full speed. None of these designs remove the LLM from steady-state fuzzing.

2.3 The off-hot-path paradigm

G²FUZZ [14] invokes the model *only* when the underlying fuzzer fails to find new coverage, and it does so by asking the LLM to synthesize a Python input generator. The generator is executed to produce a batch of candidate inputs, which AFL++ then mutates as usual. The model contributes strategy at coarse temporal granularity; mutation throughput remains native. G²FUZZ shows this paradigm is practical on 34 non-textual input formats and outperforms AFL++, Fuzztruction, and FormatFuzzer in coverage and bug-finding. We treat the off-hot-path placement as a starting *contract* that FuzzPilot inherits rather than a contribution we claim.

2.4 Why structured text needs a different intermediate

G²FUZZ’s intermediate representation is *Python code*: the LLM emits a generator script and the fuzzer executes it. Two properties of the non-textual binary setting make this choice attractive: (i) binary formats have rich constructor APIs (e.g., `PIL.Image.new` for JPEG/TIFF) that make generator scripts compact, and (ii) the generated bytes are opaque, so AFL++’s bit/byte mutators are appropriate downstream.

Structured text parsers exhibit the opposite properties. The relevant edits at the syntactic level — inserting a key, replacing a literal with an interesting value, splicing two well-formed sub-trees, applying a dictionary token — are easier to express as small *data* (an opcode, a target offset, a replacement literal) than as code. Furthermore, running LLM-emitted code at runtime carries non-trivial safety risk (network calls, file I/O, unbounded loops) that a schema-validated data structure avoids. §4 develops this contrast in detail.

2.5 Open design questions for structured text

Within the off-hot-path paradigm, three design questions remain open for structured text parsers and motivate the rest of the paper:

- **Q1: Intermediate representation.** Code-as-generator or data-as-recipe? (§4)
- **Q2: Proposal validation.** How does the controller decide whether a model proposal is worth a multi-hour main-loop commitment? (§5)
- **Q3: Static context source.** When the target’s source is unavailable, where does the structured prompt context come from? (§5.2)

3 FuzzPilot Design

3.1 Architectural contract

The architecture (Figure 1) is organized around a single invariant: *no model call appears on any code path executed by AFL++ during per-mutation work*. Model calls occur only when (a) the plateau detector marks a window as stalled, or (b) a micro-campaign has completed and a promotion




Figure 1: FuzzPilot end-to-end architecture. Five horizontal stripes separate **external tools** (Ghidra, model client), the off-hot-path **controller** (plateau detector, blackboard, agents, micro-campaign runner), persisted **artifacts** (`static_context.json`, `proposal.json`, `recipe.json`, `main_recipes/`), the native **hot path** (AFL++ runner, corpus, custom mutator, target binary), and per-run **outputs** (`fuzzer_stats`, `coverage.csv`, `events.jsonl`, audit log). Only two edges cross stripes (drawn in red): telemetry up (3) and recipe promotion down (9). No per-mutation hot-path code ever touches the model. Numbered edges 1–9 are walked through in §3 and §5.

decision is needed. The mutator never holds an open connection to the model client; it consumes recipes from a local store that contains only validated data.

This contract is shared with G²FUZZ; we adopt it as the baseline. The remainder of this section describes the three mechanisms by which FuzzPilot specializes the contract for the structured-text-parser domain.

3.2 Component overview

Plateau detector. A coverage-window monitor that flags a plateau when fewer than θ new edges accumulate over W seconds. Hyperparameters are pinned per target (§5). On a plateau the detector freezes a snapshot of the queue and writes it to the blackboard.

Blackboard. A JSON document holding the plateau snapshot, recent fuzzer_stats, and the Ghidra-derived static context for the target binary (§5.2). The blackboard is the LLM’s only view of the campaign state.

Agents. A small set of role-specialized prompt templates (coordinator, plateau-diagnosis, scheduler, cmp-hint, dictionary, format, mutator, corpus, crash-triage). Each agent reads the blackboard and proposes a recipe or a sub-recipe; the coordinator assembles a final proposal. We do not claim multi-agent specialization as a novelty in this preprint; the role-split is an engineering convenience and is ablated jointly with the other control-path components.

Recipe store. A versioned, schema-validated collection of recipes proposed, evaluated, and (in some cases) promoted. The custom mutator reads from this store at recipe-load time only.

Micro-campaign runner. Launches an isolated N -second AFL++ run with a candidate recipe installed, computes a reward over its delta in edges and paths, and reports back. The micro-campaign shares the target binary with the main run but uses a private working directory and a snapshot of the corpus.

Custom mutator. A native AFL++ mutator that dispatches on the seven recipe operations (BitFlip, OverwriteRange, InsertToken, Arith, Splice, DeleteBlock, DictionaryOverwrite; the full enumeration is in §4.2). The mutator is the only on-the-hot-path component that touches recipes; its cost profile is reported in §6 (RQ2, F5).

4 Recipe Abstraction: Data, Not Code

4.1 Recipe schema

A recipe carries two layers of state that the implementation keeps in lock-step. The high-level form, `SeedMutationStrategy` (`include/fuzzpilot/mutation/strategy.hpp:38--49`), is what the agent emits and the controller serializes; it is lowered to a compact binary form, `CompactRecipe` (`mutators/fuzzpilot/recipe_cache.hpp:45--58`), before the custom mutator loads it. The two share the following core fields:

- **id:** stable identifier; lets repeat recipes deduplicate across plateaus and across runs.
- **selector:** which subset of corpus elements the recipe applies to. A selector is a $(mode, key)$ pair where *mode* is one of `mode` / `seed_id` / `seed_hash` / `family` and *key* is the matching string.
- **priority** and **ttl_sec:** scheduling weight and time-to-live for the recipe store.

- **operator_weights** (in compact form: **weights**): a probability distribution over the fixed 7-operation vocabulary (§4.2). The mutator samples *one* op per call by these weights; the recipe is not an ordered program.
- **focus_ranges** / **protect_ranges**: half-open byte ranges $[a, b)$ where the mutator is preferred (respectively forbidden) to write. Derived from agent inference and from the Ghidra-extracted **cmp_constants** and **branch_constraints** (§5.2).
- **dictionary_tokens** (compact: **tokens**): literal byte sequences that the **InsertToken** and **DictionaryOverwrite** ops draw from. Populated from Ghidra **magic_tokens** and from agent-proposed tokens.
- **expected_signal**: the agent’s prediction of what kind of coverage gain this recipe should produce. Used by the audit log only; never consulted by the mutator.
- **fields** (compact only): per-field overrides for structured parsers, e.g. “JSON object values” or “PNG IHDR block”.

Recipes that fail JSON-schema validation are discarded before they reach the recipe store. The mutator never parses free-form model output — it only loads **CompactRecipe** entries from a local binary cache.

4.2 Operation vocabulary

The op vocabulary is small (seven ops) and closed (**mutators/fuzzpilot/recipe_cache.hpp:9--17**):

- **BitFlip** — single-bit flip at a chosen offset.
- **OverwriteRange** — write random bytes into a contiguous range.
- **InsertToken** — splice a token from **dictionary_tokens** at a chosen offset.
- **Arith** — arithmetic on a numeric field (add / subtract / multiply / divide).
- **Splice** — replace a sub-range with bytes lifted from another queue element.
- **DeleteBlock** — excise a contiguous range.
- **DictionaryOverwrite** — overwrite bytes at a chosen offset with a dictionary token.

Each call to the mutator’s fuzz function invokes **recipe.choose_operator(rng)** to sample one op by **operator_weights** and dispatches it through a 7-arm switch (**mutators/fuzzpilot/mutator_core.cpp:292--318**). There is no opcode interpreter and no nested control flow: the mutator’s work per mutation is bounded by the cost of the chosen op, which is the same order of magnitude as AFL++’s native havoc operators. This keeps the throughput cost of recipe dispatch within the budget we report in §6.3.

4.3 Example recipe (illustrative)

Listing 1 shows a recipe an agent might emit after inspecting a JSON parser’s Ghidra context (an opening brace **0x7b** expected at offset 0; **strcpy** call site detected at offset **0x108**). The mutator loads the recipe from cache and, on each call, samples an op by **operator_weights** and applies it inside **focus_ranges**, drawing tokens from **dictionary_tokens** when an **InsertToken** or **DictionaryOverwrite** op is chosen. The agent never re-enters the loop during mutation.

Listing 1: Example recipe (cJSON, focus on object/array boundary tokens).

```

1 {
2   "id": "p1_e1b_r03_plateau_07",
3   "selector": {"mode": "seed_hash", "key": "9c4f...a7b1"},
4   "priority": 3,
5   "ttl_sec": 1800,
6   "operator_weights": {
7     "InsertToken": 0.35,
8     "DictionaryOverwrite": 0.25,
9     "Splice": 0.15,
10    "OverwriteRange": 0.10,
11    "BitFlip": 0.10,
12    "Arith": 0.03,
13    "DeleteBlock": 0.02
14  },
15  "focus_ranges": [[0, 1], [42, 64]],
16  "protect_ranges": [[16, 20]],
17  "dictionary_tokens": ["{", "}", "[", "]", "\"", "true", "null"],
18  "expected_signal": "exercise object/array nesting boundary"
19 }

```




Figure 2: Recipe lifecycle. Each LLM proposal is first schema-validated, then evaluated in an N -second isolated micro-campaign, and only the top- K by reward are promoted to the main run. Discarded recipes are kept in the audit log.

4.4 Recipe lifecycle

The lifecycle (Figure 2) is: (1) **Proposal** from one or more agents; (2) **Schema validation** (discard malformed proposals); (3) **Micro-campaign** isolated N -second AFL++ run with the proposal installed; (4) **Reward** computed over Δ_{edges} and Δ_{paths} ; (5) **Promotion** of top- K proposals to the main recipe store, others archived to the audit log.

4.5 Why data, not code? Contrast with code-as-generator

G²FUZZ’s input-generator approach is the natural intermediate for non-textual formats, but in the structured-text setting data-as-recipe has three concrete advantages.

Runtime safety. Executing LLM-emitted Python at fuzz time exposes the harness to a class of risks that schema-validated data does not: import side effects, file I/O, network calls, and unbounded loops are all reachable from a free-form generator. The recipe dispatcher has no I/O capability beyond reading bytes from the mutated buffer and writing bytes back.

Cacheability and version control. A recipe is a small (≤ 2 KB) JSON document with a stable schema; thousands of recipes can be diffed, deduplicated, indexed by hash, and re-evaluated without re-prompting the model. A generator script is larger, depends on Python and library

versions, and is awkward to de-duplicate or diff.

Schema-validatable in finite time. The recipe schema is a closed grammar; validation is a JSON-schema check. A generator script is Turing-complete and cannot be statically validated for behavior; G²FUZZ relies on a syntactic-validator plus a sandbox to compensate.

We do not claim data-as-recipe is universally superior; in the non-textual setting it is plausibly less expressive. The claim is narrower: for structured text parsers, data-as-recipe yields a safer and more reusable intermediate than code-as-generator, and our evaluation (§6) reports its concrete cost (fp-active vs fp-empty).

5 Plateau Detection, Ghidra Context, and Micro-Campaign

5.1 Plateau detection

The plateau detector inspects a sliding window of W seconds and flags a plateau when, within that window, both (a) fewer than θ_e new executions and (b) fewer than θ_p new paths have accumulated. The three values are pinned per target in the run configuration (`experiments/targets/cjson/config.yaml`); for cJSON we use $W = 10$ s, $\theta_e = 50$ execs, $\theta_p = 1$ path. In each of the 5 `full-agent` runs the detector fired exactly one plateau cycle within the 14,400s budget (corroborated by the `corpus_snapshot` event count of 1 per run in `events.jsonl`); cJSON is small enough that one plateau is sufficient to use the productive-then-plateau window once. On a flagged plateau the detector freezes a queue snapshot, writes it to the blackboard, and signals the coordinator agent.

5.2 Ghidra-derived static context

Extraction pipeline. Once per target binary, FuzzPilot launches Ghidra `analyzeHeadless` in an ephemeral project directory (`<run_dir>/ghidra_project`, automatically deleted on exit) and runs a custom post-script, `FuzzPilotGhidraExtract.java`. The script writes `static_context.json` with the following schema:

- `functions[]`: top-20 risky functions, each annotated with name, address, size, list of risky calls, and a danger score. The score weights direct calls to known unsafe primitives: $10 \times \#\{\text{gets, system, popen}\} + 5 \times \#\{\text{strcpy, strcat}\} + 4 \times \#\{\text{sprintf}\} + 3 \times \#\{\text{malloc, free, memcpy, memmove}\} + 2 \times \#\{\text{printf}\}$.
- `magic_tokens[]`: up to 150 candidate dictionary entries — defined strings from `getDefinedData()` and 4-byte ASCII constants, in both endiannesses.
- `cmp_constants[]`: up to 100 scalar operands appearing in comparison instructions, useful for the `numeric_interesting` op of the recipe schema.
- `branch_constraints[]`: up to 150 tuples of the form (address, condition mnemonic, comparison mnemonic, operand 1, operand 2), recovered by a 4-instruction lookahead preceding each conditional jump.
- `decompiled_logic{}`: the decompiled C source for the top-5 risky functions, truncated to 4 KB each, for prompt-readable context.
- `structs`: target-defined struct layouts when available.

Two consumers, one extraction. The same `static_context.json` feeds two downstream components without re-entering the model:

1. **Agent blackboard.** The JSON is merged into the blackboard’s `static_analysis_context` field and read by every agent prompt during plateau handling. The LLM sees the ranked risky functions and decompiled C as natural-language context, alongside the queue snapshot and recent telemetry.
2. **Micro-campaign AFL dictionary.** The `magic.tokens` and `cmp_constants` are also converted into an AFL-format dictionary file (`static_generated.dict`) that the micro-campaign passes to AFL++ via `-x`. This puts static-analysis-derived tokens directly in the hands of the in-loop mutator, with no per-mutation LLM call.

This dual-consumer design is a deliberate choice: the LLM uses the context for high-level proposal, while AFL++ uses it for low-level token-level mutation, and both share a single, cacheable extraction.

Caching. The first run of a campaign emits `<run_dir>/static_context.json` and, if no prior cache exists for the binary, renames it to `base_intelligence.json`. Subsequent runs on the same binary load the cached JSON rather than re-running Ghidra. A 4-KB decompilation cap and the top-20 / top-150 / top-100 limits keep the artifact under 1 MB even for large targets.

Failure handling. If `analyzeHeadless` exits non-zero, FuzzPilot writes an error envelope into `static_context.json` rather than an empty object. Downstream code sees the error explicitly, and the audit log records the failure so a re-run can replay the same condition. The plateau handler still proceeds (with an empty static context); the LLM’s proposals are noted as “static context unavailable.”

Why Ghidra (and an honest caveat). We choose Ghidra over LLVM IR or angr for three operational reasons: (a) Ghidra accepts stripped binaries without source code or rebuild flags, broadening applicability; (b) its decompiled C output is closer to what a human reverse engineer would feed an LLM, and our prompts are trained on that style; (c) Ghidra headless is a stable production tool with predictable resource use, suitable for unattended fuzz campaigns. We treat the channel as an **enabling component** of the architecture, not as a free-standing novelty: on our open-source targets, an LLVM-IR or angr-based pipeline could in principle supply the same fields. Ablation E2c (§6.5) quantifies the contribution of the entire channel to coverage, jointly measuring both consumers (agent blackboard and AFL dictionary); it does not isolate Ghidra-versus-alternative-backend gains.

Channel yield per target (Table 1). The Ghidra channel’s empirical value depends on how much structurally meaningful token material the target binary actually exposes. As a sanity check on this axis, we extract the `.rodata` string-literal vocabulary from each target binary via `objcopy --only-section=.rodata` followed by GNU `strings` and a printable-ASCII filter (`scripts/paper01/extract_strings_dict.py`). This is a cheaper, library-free proxy for the broader Ghidra extraction pipeline (which the script also wraps in `extract_ghidra_dict.sh`); we use it here as the headline “how big is the candidate dictionary” number because it is deterministic, language-agnostic, and easily reproducible without a Ghidra install. Table 1 reports the count for each binary checked into the repository. *cJSON* yields only 40 usable string literals — mostly JSON

Target	Binary size	Unique tokens	Token examples (top by occurrence)
cJSON	228 KB	40	<code>null, true, false, %1.15g, u%04x</code>
libxml2 (parser)	5.6 MB	2,572	<code>"Start tag expected", "PEReference forbidden", "DOCTYPE improperly ter</code>
sqlite3 (parser/VM)	7.2 MB	1,164	<code>"SQLite format 3", ":memory:", "(join-%u)", "(subquery-%u)", unix-none</code>
openssl X.509	12 MB	6,656 [†]	ASN.1 OID labels + embedded test certificate blobs (mixed quality; see caveat)

Table 1: Per-target `.rodata` string-literal yield, on the binaries checked into `experiments/targets/`.

[†]openssl’s count is inflated by embedded test certificates and base64 blobs that survive the printable-ASCII filter but are not semantically useful as fuzzing tokens; the effective usable count is smaller. The cJSON row anchors the preprint null result (the static channel has only ~ 40 things to say to the LLM); the venue rows show that the channel has 30–160 $\times$ more raw material to work with on richer binaries. Reproduce with `python3 scripts/paper01/extract_strings_dict.py <binary>`.

keywords (`null, true, false`) and format specifiers (`%1.15g, u%04x`) — which is one structural reason the LLM proposal layer produced a null result on cJSON (§6.4): the static channel simply has very little to work with. The venue targets yield 1,000+ to 6,000+ tokens (XML diagnostic strings, SQL keywords, ASN.1 OID labels), which gives the LLM enough material to propose target-specific recipes rather than rely on the fixed-three-token DictionaryAgent fallback. We caution that *the raw token count is an upper bound on usable dictionary candidates*: openssl in particular includes embedded test PEM certificates and base64 blobs that survive the printable-ASCII filter but are not semantically useful as fuzzer dictionary entries; downstream consumers (AFL++ `-x` and the FuzzPilot mutator) re-rank by per-token reward signal.

5.3 Micro-campaign protocol

Each plateau cycle launches a fixed bundle of K_{cand} candidate recipes (default $K_{\text{cand}} = 4$, set in `micro_campaign.num_candidates`), one per intervention type (default, `dictionary, seed_focus, per_seed_recipe`). Each candidate is evaluated as follows:

1. Snapshot the current corpus and queue state to a private working directory.
2. Launch AFL++ for N seconds (default $N = 20$, set in `micro_campaign.budget_sec`) with the candidate recipe installed in the recipe store. The micro-campaign uses an edge-coverage map disjoint from the main run.
3. Compute reward

$$R = \alpha \Delta_{\text{edges}} + \gamma \Delta_{\text{crashes}} + \delta_h h - \delta_m m,$$

where h and m are recipe-hit and recipe-miss counts; if the AFL bitmap is unavailable, $\beta \Delta_{\text{paths}}$ substitutes for $\alpha \Delta_{\text{edges}}$. The weights are pinned in `src/micro/evaluator.cpp`: $\alpha=1.0$, $\gamma=10$, $\delta_h=10^{-3}$, $\delta_m=5 \times 10^{-4}$, $\beta=0.5$ (fallback only).

4. Rank candidates by R and promote the top scorer to the main recipe store *only if* $R > 0$; otherwise emit `winner_decided:status=no_significance` and `promotion_skipped:reason=no_successful_mic`, leaving the main loop on the controller’s default rule recipe.

The micro-campaign trades a small amount of main-run budget (typically N seconds at the start of each plateau-triggered cycle, $\leq 5\%$ of total wall-clock) for empirical evidence that a proposal helps. Without the gate, every model proposal would race into the main loop; bad proposals would degrade the campaign for an indeterminate time before being noticed via aggregate metrics. Ablation E2b

(`no-mutator`, controller + LLM but AFL++’s default mutator instead of the recipe-guided one) and E2c (`no-static-analysis`, LLM without Ghidra binary context) in §6 quantify the cost of trusting the model directly or removing the static channel.

5.4 Determinism and reproducibility

We pin the model temperature to 0.0 and report per-run schema-validity and fallback rates instead of relying on identical model output across runs. Each step (blackboard → proposal → recipe → reward → promotion) is written to an append-only audit log, indexed by proposal hash. The full decision trail of any run is reproducible end-to-end from the audit log plus the pinned corpus snapshot, even when the underlying model is itself replaced.

6 Preliminary Evaluation

6.1 Setup

Hardware and OS. All AFL++ campaigns and microbenchmarks run on a single 4-core Intel Xeon E5-2699 v4 @ 2.20 GHz virtual machine with 8 GB RAM, Ubuntu 24.04 LTS, kernel 6.8.0-48-generic, on a dedicated 80 GB data volume. `/proc/sys/kernel/core_pattern` is set to `core` (default on this host is an apport pipe that AFL++ refuses to run under).

Toolchain. AFL++ v4.21c built from source (pinned commit; instrumentation via `afl-clang-fast`). Ghidra 11.1.2-PUBLIC headless (OpenJDK 17). FuzzPilot at commit 85223d8 on branch `main`; Docker image digest pinned in the reproducibility appendix. Python deps (pandas 2.1.4, matplotlib 3.6.3, PyYAML 6.0.1) from Ubuntu packages.

Target. The sole evaluated target is cJSON (`experiments/targets/cjson/cjson_fuzzer`, ELF 64-bit, `afl-clang-fast` instrumented, `AFL_MAP_SIZE=65536`), seeded from 26 hand-curated JSON files covering objects, arrays, nested structures, Unicode escapes, numeric edge cases, and a few intentionally malformed variants. The repository also contains a libpng harness (`libpng_fuzzer`, `AFL_MAP_SIZE=131072`, 21 seed PNGs) but its 4-hour campaigns are not part of this preprint and are deferred to the venue version.

Modes and repeats. Five modes share a single AFL++ binary and target build, differing only in the controller’s `--ablation` switch: `baseline-afl` (no FuzzPilot controller, no model, no mutator), `full-agent` (all subsystems enabled), `rule-only` (controller + native mutator, LLM disabled), `no-mutator` (controller + LLM, AFL++ default mutator), `no-static-analysis` (controller + LLM + mutator, Ghidra context muted). Repeats per cell: 5 for `baseline-afl` and `full-agent`, 3 for the three ablation rows.

Budget. Each AFL++ run has `main.budget_sec = 14,400` (4 h wall-clock). Acceptance gates per run, from `experiments/manifests/paper01_preprint.yaml`: `run_time ≥ 14,000 s`, `execs_done ≥ 106`, `coverage.csv ≥ 200 rows`, all required artifacts present (`fuzzer_stats`, `coverage.csv`, `events.jsonl`, `run_metadata.json`, `report.md`), and `status == completed`.

Component	Pinned value
Host CPU	4-core Intel Xeon E5-2699 v4 @ 2.20 GHz
Host RAM	8 GB
OS / kernel	Ubuntu 24.04 LTS / Linux 6.8.0-48-generic
Data volume	80 GB ext4 (/dev/vdb1 on /www)
core_pattern	core (default /usr/share/apport/... overridden)
AFL++	v4.21c, built from source, afl-clang-fast instrumentation
Ghidra	11.1.2.PUBLIC headless, OpenJDK 17
FuzzPilot tag	85223d8 on main
AFL_MAP_SIZE	65536 (cJSON)
Python deps	pandas 2.1.4, matplotlib 3.6.3, PyYAML 6.0.1 (Ubuntu pkgs)
Reference LLM	deepseek-chat, temperature 0, api.deepseek.com
Fallback LLM	meta/llama-4-maverick-17b-128e-instruct via NVIDIA NIM
Per-run budget	14,400 s (4 h) wall-clock
Repetitions	5 (baseline-afl, full-agent) / 3 (ablations)

Table 2: Pinned evaluation setup. All numbers in this paper come from this single configuration. The reproducibility appendix records the Docker image digest, target binary SHA-256, and seed-corpus tarball checksum.

Model. The reference model for `full-agent` runs is `deepseek-chat` at temperature 0 via the OpenAI-compatible endpoint <https://api.deepseek.com/chat/completions>. A secondary configuration uses `meta/llama-4-maverick-17b-128e-instruct` via NVIDIA NIM (<https://integrate.api.nvidia.com>); the configuration switch is documented in the reproducibility appendix. The microbench (E3, §6.3) does not invoke any LLM.

Parallelism and isolation. AFL campaigns run 4-way parallel (one worker per core); microbench runs (E3) are run serially with *all AFL workers suspended via SIGSTOP* for the duration of the microbench (≈ 30 s; resumed via *SIGCONT*). Empirically, running the microbench either concurrently with AFL (CPU contention) or 4-way parallel against itself inflates per-rep variance by 2–3 $\times$ and inverts the paper’s central RQ2 ratio (§6.3); the isolation methodology is therefore a prerequisite for sound microbench numbers on a 4-core host. The ≈ 30 s SIGSTOP window consumes $< 0.2\%$ of each AFL run’s 14,400 s budget and is logged in `results/paper01_ai_recipe_mutation/runs/_logs/`.

6.2 RQ1 — Plateau recovery on cJSON (E1a vs E1b)

Edge ceiling on cJSON is a target-side property. Both modes saturate at the same 269-edge ceiling (`edges_found` field of AFL++’s `fuzzer_stats`, also corroborated by `bitmap_cvg` converging to 23.21% in all 10 runs) under the 14,400s budget: `baseline-afl` reaches 269 in 5 / 5 runs; `full-agent` reaches 269 in 4 / 5 runs and 266 in the remaining one (r02). Figure 3 plots both median curves on the same axes; they converge to the dashed ceiling line and do not exceed it in any mode. This is a property of the cJSON harness, not of the fuzzer: cJSON exposes a limited reachable-edge space that AFL++ alone exhausts well before the budget, leaving no headroom for plateau-recovery improvement at the *ceiling* level. We therefore use cJSON for RQ2 (throughput parity) and RQ3 (LLM quality), and defer ceiling-level RQ1 evidence to larger targets in the venue paper.

Plateau-duration delta is the real RQ1 signal. Even when both modes saturate to the same ceiling, the moment of saturation differs systematically. Table 3 and Figure 5 report the median `last_find` (wall-clock at which the last new edge was discovered, i.e. end of the productive phase)

Coverage on cJSON (zoomed; *first 60 s ramp 180→260+ omitted*)




Figure 3: Edge-discovery trajectory on cJSON for **baseline-afl** (E1a, $N=5$) and **full-agent** (E1b, $N=5$), zoomed into 264–270 edges. Both modes ramp rapidly from 180 to ≥ 264 in the first 60 s (omitted here for clarity) and then converge to the same 269-edge ceiling. The shape *within the zoom band* is what differs: **baseline-afl** reaches 269 at a per-run median of $\approx 2,500$ s (slowest run: 4,448 s; range [1,081, 4,448]) and remains flat, whereas **full-agent** reaches 269 at a per-run median of $\approx 6,700$ s (one run never reaches 269 within budget; range over the four runs that do: [3,487, 14,245]) but continues to discover new edges (268 then 269) much later into the budget — the median **last_find** of 13,059 s is 1,148 s after baseline’s 11,911 s. The takeaway is not that **full-agent** reaches the ceiling faster (it does not: it reaches the ceiling *later* on average), but that the productive phase is stretched out closer to the budget end, which shortens the trailing plateau. The two coloured dashed verticals mark the per-mode median **last_find**: 11,911 s for **baseline-afl** vs 13,059 s for **full-agent**.

and the resulting *plateau* duration (gap between **last_find** and the 14,400 s budget end). The full-agent group extends the productive phase by 1,148 s (≈ 19 min) at the median, corresponding to a descriptive 45.3% reduction in plateau duration ($2,532 \rightarrow 1,384$ s). *This delta is descriptive only*: with $N=5$ per arm the per-run plateau distributions [539, 1,185, 2,532, 3,314, 3,970] s and [238, 835, 1,384, 1,610, 3,226] s overlap substantially, and a two-sided Mann–Whitney U test (following [3, 10]) gives $U=8$, $p=0.42$; the Vargha–Delaney $A_{12}=0.68$ [20] indicates a moderate effect in the direction the median suggests, but the bootstrap 95% CIs for the two medians ([539, 3,970] s for baseline, [238, 3,226] s for full-agent) almost entirely overlap. We report the descriptive -45.3% as the central point estimate and defer significance testing to the venue version at $N \geq 10$ on non-saturated targets, where formal rank-sum tests have the statistical power to discriminate. One full-agent run (r04) keeps discovering new edges until 14,205 s — ≈ 3 minutes before the budget exhausts — versus the baseline best of 13,907 s (r05). The point estimate is consistent with the recipe-guided mutator re-distributing havoc-style attempts via its default **DictionaryAgent** recipe (fixed token + operator-weight distribution; see §5.3), which delays the plateau even when the same final edge set is eventually reached. We attribute the descriptive delta to the mutator framework rather than to LLM-promoted recipes: §6.4 reports that none of the 20 LLM candidates evaluated by the micro-campaign gate across the 5 **full-agent** runs were promoted to the main store (every candidate’s $R=0$ because cJSON has already saturated by plateau time), so the main loop in every

Mode	median <code>last_find</code> (s)	median plateau (s)	95% CI plateau	Δ plateau (descriptive)	Mann–Whitney U (p , A_{12})
baseline-afl (E1a, $N=5$)	11,911	2,532	[539, 3,970]	—	—
full-agent (E1b, $N=5$)	13,059	1,384	[238, 3,226]	−1,148 (−45.3%)	8 (0.42, 0.68)

Table 3: Plateau-duration delta on cJSON. *Both medians and the absolute delta are descriptive only:* the two distributions overlap heavily, the Mann–Whitney p -value (one-sided test against the null “baseline plateau = full-agent plateau”) is 0.42, and the bootstrap 95% CIs for the medians overlap. The Vargha–Delaney A_{12} effect size of 0.68 indicates that in $\approx 68\%$ of random pairs the baseline-afl plateau is longer than the full-agent plateau — a moderate but not decisive effect at $N=5$. Raw per-run values are in Figure 5.

E1b run ran the same default recipe that the **rule-only** ablation E2a (§6.5) runs on. A pure measurement of “mutator framework vs vanilla AFL havoc” is therefore E2a vs E1a; E1b vs E1a adds the LLM and micro-campaign *infrastructure cost*, which Table 3 bounds at well under the observed plateau gain. *However, as §6.5 discusses, the ablation ordering (no-mutator median plateau 930 s < rule-only 1,165 s < full-agent 1,384 s) contradicts the naive “mutator-framework drives the gain” attribution and points toward the controller’s plateau detector and corpus-snapshot reorganization as a co-driver; the attribution remains a descriptive hypothesis at this N .*

Per-run timeline. Figure 5 shows the full per-run picture. Three full-agent runs (r01, r03, r04) push past the *baseline median* (11,911s) and stay productive significantly longer; one (r04) is productive until 14,205s. Two full-agent runs (r02, r05) plateau earlier than the baseline best (13,907s); this is expected on a saturated target, where which corpus subset a worker ends up in dominates over agent quality. With $N=5$ reps the per-mode signal is qualitative; the venue paper will repeat this on Magma-suite targets at $N \geq 10$ with formal Mann–Whitney U and Vargha–Delaney A_{12} reporting, following Klees et al. [10]’s evaluation guidelines.

Statistical methodology disclosure. For the descriptive comparisons above and in §6.5 we use the two-sided Mann–Whitney U test (`scipy.stats.mannwhitneyu`, exact for $N \leq 8$) for non-parametric difference in medians, the Vargha–Delaney A_{12} [20] for effect size, and the percentile bootstrap (10,000 resamples) for median confidence intervals. The code that produces all reported numbers (`scripts/paper01/review_stats.py`) is in the repository and run from the same per-run `fuzzer_stats` files cited in the reproducibility appendix (§9). Per Klees et al. [10], the canonical evaluation guideline calls for $N \geq 20$ repetitions; *this preprint’s $N=5$ / $N=3$ does not meet that bar* and we therefore present all comparisons as descriptive point estimates with their non-parametric companions, rather than as significance claims.

Time-to- N -edges (Klees-style endpoint-free metric). Endpoint metrics like plateau duration depend only on `last_find` and ignore the shape of the discovery curve; Klees et al. [10] §6.1 caution against relying on them. We therefore report Table 4: the wall-clock at which each run first reaches N edges, for $N \in \{264, 268, 269\}$ (264 is one edge below the steady-state ramp, 268 is the penultimate edge, 269 is the ceiling).

The complementary endpoint-free reading is that the *ramp-up* (time to 264) is mode-independent; the *tail* (time to 269) is where modes diverge, and **baseline-afl** actually reaches the ceiling *faster* at the median than **full-agent**. **full-agent**’s advantage in plateau duration (Table 3) therefore comes from discovering its final edge *later in the budget*, not from discovering the ceiling earlier.

Mode	median time to 264 edges (s)	median time to 268 edges (s)	median time to 269 edges (s)
baseline-afl ($N=5$)	60	660	2,524
full-agent ($N=5$)	120	662	6,702 [†]
rule-only ($N=3$)	60	661	1,923
no-mutator ($N=3$)	60	601	1,804

Table 4: Median per-run wall-clock to reach N edges. [†]One **full-agent** run (r02) never reaches 269 within budget; the median is over the four runs that do {3,487, 4,329, 9,076, 14,245} s. All four ablation modes reach 264 and 268 in roughly the same time as **baseline-afl**; the gap opens at the 269 ceiling, where **baseline-afl** arrives *earlier* on average than **full-agent** but then plateaus longer.

The two readings — “plateau is shorter” and “ceiling is reached later” — are internally consistent: a 14,400s budget that ends with a new edge at 14,205s (full-agent r04) has a tiny plateau but a very late ceiling-touch. A non-saturated target would let the two metrics diverge in informative directions.

Mutation-rate redistribution (per-cycle vs per-execution). A second consequence of the recipe-guided mutator is visible in `cycles_done` (Figure 4). The **baseline-afl** median is 797 corpus cycles inside the 14,400s budget; the **full-agent** median is 140, and the two ablations that engage the same mutator (**rule-only**: 119; **full-agent**: 140) sit in that same $\approx 5\text{--}7\times$ -fewer range, whereas **no-mutator** (671) tracks much closer to **baseline-afl**. The ratio is therefore attributable to the recipe-guided mutator itself, not to any LLM-introduced step. With a comparable total execution count (within 6%; RQ2, Figure 6), full-agent spends $\approx 5.7\times$ more executions on each pass through the queue before recycling to the top — exactly the mutation-rate redistribution the default recipe encodes (*focus dictionary-token insertions and overwrites over the first 4 KB*, with fixed operator weights pinned at `insert_token=0.35` / `dictionary_overwrite=0.25` / `overwrite_range=0.20` / `arith=0.10` / `bit_flip=0.10`). *We caution that `cycles_done` is not a cross-mode comparator for “progress”*: it is a derived quantity of the form `execs_done/execs-per-cycle`, where the denominator depends on per-input fuzzing intensity and not on coverage progress. Two modes can run the same total executions and reach the same edge ceiling with very different `cycles_done` values; the apparent “ $5.7\times$ less work per cycle” interpretation reduces to “the recipe-guided mutator runs more per-input mutation steps before popping the queue head”, which is a definitional property of the recipe and not a finding. Whether the concentration is empirically beneficial is a target-specific question; on cJSON it correlates with the longer productive phase (Figure 5) but not with a larger final edge set, because the target ceiling at 269 binds.

Per-run summary table (T1). Table 5 reports the 10 completed runs (5 baseline + 5 full-agent) at end-of-budget. Acceptance gates pass on all 10 (`run_time` $\geq 14,000$ s, `execs_done` $\geq 10^6$); zero saved crashes in either mode (expected for cJSON’s safety-tested code base).

Extended per-run statistics (T1-ext). Table 6 reports the time-budget breakdown, target stability, and plateau metrics for all 10 completed runs (5 baseline + 5 full-agent). The fuzz-time fraction (`fuzz_time/run_time`) is uniformly $\geq 99.7\%$ in both modes, indicating that calibration, trim, and sync overheads — as well as the ≈ 30 s SIGSTOP windows from the E3 microbench (§6.3) — did not noticeably perturb the AFL campaigns. Target stability (AFL++’s coverage-bitmap




Figure 4: `cycles_done` per run across the two modes ($N=5$ each). Solid horizontal bars are per-mode medians. `full-agent` completes $\approx 5.7\times$ fewer corpus cycles in the same 14,400 s budget, despite an *equal-or-higher* total execution count (Figure 6): the recipe-guided mutator spends more work per corpus item before recycling.

stability check) is $\geq 99.25\%$ across all 10 runs, well within the standard “stable target” threshold of 90%, so all measurements are comparable across modes.

Cross-run throughput consistency. Within `baseline-afl`, the 5-run `execs_per_sec` spread is $\{12,826, 12,958, 13,049, 13,108, 13,854\}$, $CV = 3.1\%$ (median 13,049). Within `full-agent`, the 5-run spread is $\{13,540, 13,542, 13,816, 13,949, 15,597\}$, $CV = 6.1\%$ (median 13,816). Both groups satisfy the `execs_done` $\geq 10^6$ and `run_time` $\geq 14,000$ acceptance gates with $> 10^3\times$ and $\sim 1.03\times$ margins respectively. The end-to-end RQ2 throughput-parity ratio (§6.3) is computed from these numbers.

6.3 RQ2 — Mutator throughput parity (F5 microbench)

RQ2 has two pieces of evidence. End-to-end `execs_per_sec` parity between `baseline-afl` (E1a) and `full-agent` (E1b), measured by the runtime `fuzzer_stats:execs_per_sec` field over the 14,400 s budget, is the canonical RQ2 number; Table 7 reports it. The microbench (E3) measures the dispatch cost of the recipe-guided mutator alone (no AFL, no target); the rest of this subsection covers the microbench.

End-to-end throughput parity (RQ2 part a). Across $N=5$ runs per mode, the median `execs_per_sec` ratio $\text{median}(E1b)/\text{median}(E1a) = 13,816/13,049 = 1.059\times$ (Figure 6). The manifest acceptance gate requires this ratio to be ≥ 0.85 for throughput parity; the observed ratio is $1.059\times$, which is not only *within* the gate but exceeds it by a wide margin: `full-agent` is on average 5.9% *faster* than baseline despite the recipe-guided mutator and the agent-runtime overhead.

Why is full-agent not slower? Three structural reasons. First, the recipe-guided mutator’s dispatch is a 7-op switch table (§4.2), which is structurally simpler than AFL++’s open-ended havoc selection (see also F5 microbench below). Second, the LLM is off the hot path by design (§3.1); it is invoked at plateau boundaries only, not per mutation. Third, the micro-campaign consumes $K_{\text{cand}} \cdot N \approx 80$ s of main-run wall-clock per plateau ($< 0.6\%$ of the 14,400 s budget), which is small enough not to register in the end-to-end `execs_per_sec` numbers. The end-to-end measurement folds all three effects into a single number.




Figure 5: Per-run plateau-onset timeline for E1a (top, blue) and E1b (bottom, orange). Solid coloured bars = productive fuzzing (zero or more new-edge discoveries per AFL++ tick within the bar; the bar ends at the run’s last new-edge timestamp, `last_find`). Hatched red bars = the trailing plateau interval (zero new-edge discoveries between `last_find` and budget end). The blue / orange vertical dashed lines mark the per-mode median `last_find` (11,911s vs 13,059s). `full-agent` extends the median productive phase by ≈ 19 min (-45.3% plateau duration).

Per-run breakdown. Within E1b, the per-rep `execs_per_sec` values are $\{13,540, 13,542, 13,816, 13,949, 15,597\}$, all above the lowest baseline rep (12,826). The acceptance ratio holds not just at the median but at every individual rep: even the slowest E1b rep (13,540) divided by the fastest E1a rep (13,854, r05) is $0.977\times$, still well above the 0.85 gate. We attribute r05 = 15,597 in E1b to the fact that this run was the lone occupant of the 4-core host during the second parallelism wave of the E1b campaign, identical to the way r05 in E1a (13,854) ran alone in its own wave.

Microbench setup (E3). The microbench (E3) is the off-loop dispatch-cost number: it isolates the cost of FuzzPilot’s mutator dispatcher from the cost of running AFL++ itself. All three configurations load a custom mutator `.so` through the same `dlopen` path that AFL++ uses in real runs:

- **vanilla** — `libvanilla_havoc.so`, a hand-written shim that performs AFL++ havoc-equivalent random byte/bit edits via the same custom-mutator API surface (`tools/mutator_microbench/vanilla_havoc`). This is a *cost-symmetric dispatch shim*, not a measurement of AFL++’s real internal `havoc_mutate`: both `vanilla` and the FuzzPilot configurations pay the same `dlopen` + symbol-resolution + per-call API dispatch cost, so the difference isolates the dispatcher itself

Mode	Run	run_time	execs_done	exec/s	cycles	corpus	edges	bitmap_cvg
baseline-afl	r01	14,444	188,492,319	13,049	797	585	269	23.21%
	r02	14,443	187,158,013	12,958	800	607	269	23.21%
	r03	14,443	189,334,714	13,108	702	606	269	23.21%
	r04	14,443	185,251,854	12,826	569	649	269	23.21%
	r05	14,446	200,147,594	13,854	819	607	269	23.21%
	<i>median</i>	14,443	188,492,319	13,049	797	607	269	23.21%
full-agent	r01	14,441	199,520,607	13,816	122	598	269	23.21%
	r02	14,442	195,572,895	13,542	141	573	266	22.95%
	r03	14,443	195,567,575	13,540	144	601	269	23.21%
	r04	14,443	201,466,942	13,949	126	604	269	23.21%
	r05	14,437	225,180,431	15,597	140	600	269	23.21%
	<i>median</i>	14,442	199,520,607	13,816	140	600	269	23.21%

Table 5: T1: per-run end-of-budget summary across the 10 completed E1a + E1b runs on cJSON. **full-agent** median **exec/s** (13,816) is *above* baseline (13,049); see RQ2 (§6.3). All 10 runs reach $\geq 14,400$ s and $\geq 1.85 \times 10^8$ execs. Preliminary E2a (**rule-only**) and E2b (**no-mutator**) ablation numbers ($N=3$ each) are reported in §6.5; E2c (**no-static-analysis**) reported in §6.5.

rather than AFL’s full havoc engine. A real AFL++ havoc comparator would also require running AFL++’s scheduler and the corresponding bookkeeping, which is out of scope for E3 and instead surfaced end-to-end via E1a vs E1b.

- **fp-empty** — the FuzzPilot mutator with an empty recipe store. It pays the full FuzzPilot dispatch cost (recipe cache lookup, op selection, telemetry ring write) but falls through to a vanilla-equivalent op on every call. Isolates the “FuzzPilot machinery is loaded but no recipe is active” cost.
- **fp-active** — the FuzzPilot mutator with a populated recipe (**op.weights** non-uniform, **dictionary_tokens** populated, **focus_ranges** pinned). This is the on-paper steady state during **full-agent** runs.

Each configuration runs 100,000 mutation calls over a fixed 10,000-element cJSON seed corpus, repeated $N = 5$ times. All microbench reps are run *serially* with all AFL workers **SIGSTOP**-suspended (see §6.1); resuming AFL after the 5-rep batch consumes ≈ 30 s of wall-clock, well under the 14,400 s budget per AFL run.

Primary RQ2 claim (fp-active vs fp-empty: descriptive, not formally equivalent at $N=5$). The median **fp-active**/**fp-empty** ratio is $3,008,030/2,961,850 = 1.016\times$ (inside the $[0.95, 1.05]$ band stated in §3); the *mean* ratio, however, is $2,800,142/3,222,108 = 0.869\times$ (-13.1%). The mean-vs-median divergence reflects the high coefficient of variation ($\sim 25\text{--}30\%$) at $N=5$ and the presence of an **fp-active** run (r04) at 1.7 Mexec/s alongside an **fp-empty** run (r05) at 4.4 Mexec/s. A two-sided Welch *t*-test of **fp-active** vs **fp-empty** gives $t=-0.81$, $df\approx 7.2$, $p=0.45$ — no significant difference is detected, but at $N=5$ this is a power-limited null finding rather than evidence of equivalence. *We additionally ran a two one-sided test (TOST) procedure [18] for equivalence at the $\pm 5\%$ band claimed in §3: $p_{TOST}=0.68$, so the test does **not** establish equivalence at the conventional $\alpha=0.05$ level.* The widest band the data *does* support at $\alpha=0.05$ is roughly $\pm 50\%$ ($p_{TOST}=0.028$), which is too wide to be a useful no-overhead claim. The widest band we tested ($\pm 50\%$) also does not pass TOST at this N , because the sample standard errors dominate. We therefore downgrade the original “no measurable per-call overhead” claim to: at $N=5$ on this SIGSTOP-isolated microbench,

Mode	Run	fuzz_time	run_time	fuzz %	stability	last_find (s)	plateau (s)
baseline-afl	r01	14,424	14,444	99.86%	99.26%	11,130	3,314
	r02	14,432	14,443	99.92%	99.26%	10,473	3,970
	r03	14,424	14,443	99.87%	99.26%	13,258	1,185
	r04	14,402	14,443	99.72%	99.26%	11,911	2,532
	r05	14,427	14,446	99.87%	99.26%	13,907	539
	<i>median</i>	14,424	14,443	99.87%	99.26%	11,911	2,532
full-agent	r01	14,415	14,441	99.82%	99.26%	13,606	835
	r02	14,421	14,442	99.85%	99.25%	11,216	3,226
	r03	14,425	14,443	99.88%	99.26%	13,059	1,384
	r04	14,417	14,443	99.82%	99.26%	14,205	238
	r05	14,423	14,437	99.90%	99.26%	12,827	1,610
	<i>median</i>	14,421	14,442	99.85%	99.26%	13,059	1,384

Table 6: T1-ext: per-run `fuzz_time` budget breakdown, target stability, `last_find` wall-clock, and plateau duration across the 10 completed runs. `full-agent` median plateau duration is 45.3% shorter than `baseline-afl` median (1,384s vs 2,532s); see RQ1 (§6.2). Preliminary E2a / E2b ablation medians ($N=3$) are reported in §6.5; E2c reported there (median 234s, $N=3$).

Mode	n	mean exec/s	median exec/s	SD (CV)
baseline-afl (E1a)	5	13,159	13,049	403 (3.06%)
full-agent (E1b)	5	14,089	13,816	862 (6.12%)
ratio (E1b / E1a) median	—	—	$\approx 1.06\times$	—
ratio (E1b / E1a) mean	—	1.0707 $\times$	—	—
acceptance gate	—	—	≥ 0.85	—

Table 7: RQ2 part (a): end-to-end `execs_per_sec` parity between the two modes. Both groups satisfy the `run_time` $\geq 14,000$ s and `execs_done` $\geq 10^6$ acceptance gates from the manifest. The observed median ratio is $\approx 1.06\times$ (1.0587 to four significant figures, but at $N=5$ that level of precision is overstated; dropping the r05 solo-occupancy outlier moves the ratio to $1.05\times$, well within the parity band), which is $\sim 25\%$ above the 0.85 parity threshold. Note the `full-agent` CV is slightly higher than baseline (6.12% vs 3.06%), driven primarily by r05’s outlier (15,597 exec/s); we report median to be robust to this outlier.

we (a) do not detect a significant mean throughput difference between `fp-active` and `fp-empty` ($p=0.45$, two-sided), (b) cannot formally establish equivalence within $\pm 5\%$ ($p_{\text{TOST}}=0.68$), and (c) note that the median ratio is essentially flat ($1.016\times$) while the mean ratio is wider ($0.87\times$). A formal equivalence claim requires $N \geq 30$ on this microbench — which is straightforward to do and is folded into the venue version. The dispatch-cost result is therefore best read as “no large overhead detected” rather than “no overhead exists”; the end-to-end `execs_per_sec` comparison against AFL++’s own mutator is the E1a-vs-E1b number reported in §6.2 (median ratio $\approx 1.06\times$), which is the load-bearing parity claim for the rest of the paper.

Sanity claim (fp-empty vs vanilla within $5\times$). The median `fp-empty/vanilla` ratio is $1.529\times$; the threshold is $\leq 5\times$, so the claim holds with a wide margin. The unexpected direction — FuzzPilot’s dispatcher is *faster* than the AFL-havoc-equivalent baseline, not slower — is consistent with FuzzPilot’s closed 7-op switch table being more predictor-friendly than `vanilla_havoc.cpp`’s open-ended random op selection. We retain the $5\times$ ceiling in the paper as a conservative gate (FuzzPilot’s mutator could plausibly regress under future op additions) but report the actual $1.53\times$




Figure 6: End-to-end `execs_per_sec` per run for both modes ($N=5$ each); horizontal bars are per-mode medians. Even the slowest `full-agent` rep (13,540) exceeds every `baseline-afl` rep except `r05` (13,854, also a solo-occupancy run). The median ratio 1.059 $\times$ is well above the 0.85 acceptance gate.

Config	n	mean exec/s	SD (CV%)	median exec/s	IQR	median ns/call
<code>vanilla</code>	5	2,005,078	341,644 (17.0%)	1,937,110	200,200	516.2
<code>fp-empty</code>	5	3,222,108	958,259 (29.7%)	2,961,850	1,310,390	337.6
<code>fp-active</code>	5	2,800,142	673,566 (24.1%)	3,008,030	393,750	332.4

Table 8: F5 microbench: mutator-only `exec/sec` across the three dispatch-symmetric configurations (Figure 7). $N = 5$ reps per configuration, serial execution with AFL workers SIGSTOP-suspended to remove the 4-core host’s CPU contention from the measurement. See `results/paper01_ai_recipe_mutation/microbench/` for raw per-run JSON.

as a positive side effect of the closed vocabulary.

Variance and methodology note. Per-config CV is 17–30%, which is large for a microbench but expected on a 4-core VM with shared L3 cache and noisy-neighbor effects from the underlying hypervisor (steady-state `%st` was 4–12% during measurement). The black dots overlaid on Figure 7 are the 15 individual reps (5 per config); they show the full distribution that the mean and SD summarize. Table 9 lists the per-rep numbers for reproducibility. Two alternative measurement methodologies inflate this variance and *invert* the primary RQ2 ratio:

- Running the microbench 4-way parallel (i.e. `run_batch` default for AFL campaigns) on a 4-core host: median `fp-active`/`fp-empty` ratio collapses to 0.716 $\times$ (FAIL), driven by reps that win or lose the scheduling lottery against their siblings.
- Running the microbench serially *while AFL campaigns are active*: median ratio drifts to 1.380 $\times$ (FAIL), driven by AFL workers monopolizing CPU during long mutations and yielding briefly to the microbench between syscalls.

We archive both noisy datasets at `results/paper01_ai_recipe_mutation/microbench_archive_*/` so reviewers can reproduce the methodology comparison; the main result of this section uses only the SIGSTOP-isolated serial dataset under `microbench/`.




Figure 7: Mutator-only throughput by configuration (mean $\pm$ SD, $N=5$ reps each, SIGSTOP-isolated). The median `fp-active`/`fp-empty` ratio is $1.016\times$ (the mean ratio is $0.87\times$; see body for TOST). At $N=5$ a TOST equivalence test does *not* establish $\pm 5\%$ equivalence and the two-sided t -test does not detect a significant difference either — both are power-limited null results that the venue version will resolve at $N \geq 30$. FuzzPilot’s closed 7-op dispatcher is descriptively faster than the AFL-havoc-equivalent `vanilla` baseline (median ratio $1.53\times$), a positive side effect of the closed op vocabulary that we report alongside the conservative *within-5* $\times$ sanity gate.

Config	r01	r02	r03	r04	r05
<code>vanilla</code>	2,591,470	1,770,420	1,755,770	1,937,110	1,970,620
<code>fp-empty</code>	2,961,850	3,982,240	2,080,370	2,671,850	4,414,230
<code>fp-active</code>	2,645,520	3,008,030	3,561,610	1,746,280	3,039,270

Table 9: Per-rep `exec/sec` from the 15 SIGSTOP-isolated microbench runs that feed Figure 7 and Table 8. Each cell is the `mean_exec_per_sec` field from `results/paper01_ai_recipe_mutation/microbench/p1-e3_microbench-<config>-r<NN>.json`.

6.4 RQ3 — LLM infrastructure exercise on a saturated target (E1b)

RQ3 asks how the off-hot-path LLM machinery behaves at plateau boundaries on a target where AFL++ alone already exhausts the reachable edge space (cJSON; §6.2). Concretely: do proposals parse, do they survive the micro-campaign gate, what is their dollar cost, and is the full decision trail recoverable? Table 10 summarizes the 5 **full-agent** runs of E1b.

Schema-validation rate. Each plateau cycle fires a fixed bundle of 9 agent tasks. Across $5 \times 9 = 45$ LLM responses, 43 (95.6%) produced schema-valid recipes; 2 responses (r04, r05; both tagged `error_kind="schema_invalid"` in `agent_decisions.jsonl`) failed schema validation and were dropped without invoking the rule-only fallback path (`fallback_used` is `false` on every record). With the model pinned at temperature 0, this is consistent with the model staying inside the `schema-validated recipe (data)` contract that the prompt enforces (§4.1). The agent bundle (coordinator, plateau-diagnosis, scheduler, cmp-hint, dictionary, format, mutator, corpus, crash-triage) explains why the per-run *LLM-call* record count is 9 rather than per-mutation.

Promotion rate: zero promotions on cJSON. Each plateau cycle launches $K_{\text{cand}}=4$ candidate recipes (§5.3), one per intervention type (`default`, `dictionary`, `seed_focus`, `per_seed_recipe`). Each candidate is evaluated in a private 20s AFL++ run over a corpus snapshot taken at plateau

Run	log records	schema-valid LLM resp.	schema-invalid LLM resp.	fallback_used	micro candidates	promoted (gate res)
r01	289	9	0	0	4	0 (skipp
r02	302	9	0	0	4	0 (skipp
r03	402	9	0	0	4	0 (skipp
r04	215	8	1	0	4	0 (skipp
r05	200	8	1	0	4	0 (skipp
total	1,408	43	2	0	20	

Table 10: RQ3: LLM-call quality and micro-campaign outcome across the 5 **full-agent** runs. “Decisions logged” is the total line count of `agent_decisions.jsonl`, including non-LLM events (plateau detection, recipe load, telemetry). “Schema-valid” and “schema-invalid” count records carrying a model response, separated by whether the response parses against the recipe schema (§4.1). `fallback_used` counts plateau windows where the controller fell back to the rule-only mutator on a malformed LLM proposal: 0 across all 5 runs, because the controller silently drops schema-invalid responses without triggering the fallback path. “Micro candidates” is the number of micro-campaign evaluations per run (default $K_{\text{cand}}=4$, one per intervention type). “Promoted to main” is the number of LLM-proposed recipes that the micro-campaign gate accepted into the main recipe store; see the Promotion-rate paragraph below.

time. Across the 5 runs, the controller fired exactly one plateau cycle per run, so the gate evaluated $5 \times 4 = 20$ candidates in total. Each micro-campaign ran the full 20s and executed $\approx 3.6 \times 10^5$ inputs (visible in the `events.jsonl micro_result` records). *Every single one of the 20 candidates returned* $\Delta_{\text{edges}}=0$, $\Delta_{\text{paths}}=0$, $\Delta_{\text{crashes}}=0$, so the reward (§5.3) was $R=0$ on all 20 evaluations. The controller therefore emitted `winner_decided:status=no_significance`, `winner_reward=0.0` followed by `promotion_skipped:reason=no_successful_micro_campaign` in every run, and **no LLM-proposed recipe entered the main recipe store across the 5 full-agent runs**. The main loop spent its 14,400s on the controller’s default `DictionaryAgent` recipe (`tokens={FUZZ, MAGIC, TOKEN}`, fixed operator-weight distribution), which is exactly what the **rule-only** ablation E2a (§6.5) runs on. The result is consistent with cJSON’s saturation behaviour: baseline AFL++ alone reaches the 269-edge ceiling at a per-run median of $\approx 2,500$ s (slowest baseline run: 4,448s; §6.2), leaving the parent corpus at plateau time with no headroom for a 20s micro-campaign to find a new edge. The micro-campaign gate’s refusal to promote on cJSON therefore demonstrates only one half of its intended property — the *no-false-positive* half: when reward is identically zero, the gate does not promote. The complementary property — that the gate distinguishes a good recipe from a bad recipe when reward *varies* — is **not** exercised by this evaluation, because the saturated target gives the gate no reward signal to discriminate against. The LLM proposal layer’s positive contribution on a non-saturated target therefore remains undemonstrated by this preprint, and is explicit future work in the venue paper.

Cost. Across the 5 runs, the DeepSeek API balance dropped by ¥0.07 ($\approx$ \$0.01 USD at the FX rate of \$1 USD $\approx$ ¥7.2 quoted at submission time; DeepSeek bills in Chinese yuan / RMB): from ¥8.53 at the start of E1b to ¥8.46 at the end. That is \$0.0002 per LLM call and \$0.002 per full 4-hour campaign on the LLM-API ledger. *Compute cost is omitted from this figure*: at a representative \$0.10 USD per vCPU-hour on commodity cloud, a 4-hour 4-vCPU campaign costs $\approx$ \$1.60 USD, three orders of magnitude more than the LLM API. The headline finding is therefore that the LLM API cost is *not* the dominant cost; absolute numbers should be read with that scaling in mind, and a different LLM provider (OpenAI, Anthropic, or self-hosted) would shift the API-side figure by approximately one order of magnitude without changing the qualitative conclusion. The

pricing also benefits from DeepSeek’s prompt-cache hit on the shared blackboard prefix across the 9 agent tasks in a single plateau cycle; a cold-cache estimate (without prefix sharing) would be roughly $10\times$ larger and is still well below the compute cost. The system also supports an OpenAI-compatible drop-in via NVIDIA NIM (`meta/llama-4-maverick-17b-128e-instruct`) at zero marginal cost on the NIM free tier; see the reproducibility appendix.

Audit trail. Every plateau handling cycle in `agent_decisions.jsonl` carries the blackboard hash (`context_hash`) it was conditioned on and the `response_hash` of the model’s output; combined with the schema-validated recipe stored under the plateau directory and the `events.jsonl` `micro_result` records, any decision in any of the 5 runs can be replayed end-to-end from the on-disk artifacts alone, without re-invoking the model. The same audit trail makes the null result above auditable: a reviewer can verify that the gate’s refusal to promote any candidate is grounded in $R=0$ evidence (every `micro_result` carries Δ_{edges} , Δ_{paths} , and Δ_{crashes} in its payload) rather than a controller bug.

6.5 RQ4 — Component contribution (E1b vs E2a / E2b / E2c)

The ablation matrix is intended to attribute the `full-agent` plateau-recovery delta among three components: the recipe-guided custom mutator (E2a removes the LLM layer but keeps the mutator and default rule recipe), the controller’s orchestration plus the mutator (E2b removes the mutator but keeps the controller, LLM, and Ghidra channel), and the Ghidra static channel (E2c removes Ghidra but keeps everything else). The `rule-only` (E2a), `no-mutator` (E2b), and `no-static-analysis` (E2c) campaigns have each completed $N=3$ runs of 14,400s on the experiment host (E2c required a controller-launch retry after its first batch exited before AFL telemetry was emitted, recorded as `main_afl_exited_before_telemetry` in the per-run `events.jsonl`; the retry succeeded and its data is reported below). *All comparisons in this subsection are statistically descriptive only:* with $N=3$ per ablation and $N=5$ for the main arms, no pairwise plateau or `last_find` difference reaches the conventional $p<0.05$ threshold under a two-sided Mann–Whitney U test (see §6.2), and the bootstrap medians for the ablations overlap substantially with each other and with `full-agent`.

E2a rule-only (preliminary, $N=3$). With the LLM proposal layer disabled and the controller emitting its default `DictionaryAgent` recipe directly, the 3 `rule-only` runs reach the same 269-edge ceiling that `baseline-afl` and `full-agent` reach, with median `last_find` = 13,309s and median plateau = 1,165s (descriptive -54.0% vs the baseline median of 2,532s; Mann–Whitney $U=4$, $p=0.39$; Vargha–Delaney $A_{12}=0.73$). Per-run plateaus are [674, 1,165, 1,540]s and the bootstrap 95% CI for the median is [674, 1,540]s. Median `execs_per_sec` is 13,514 and median `cycles_done` is 119. The point estimate suggests E2a is slightly faster than `full-agent` (median plateau 1,165s vs 1,384s), but the per-run ranges [674, 1,540]s and [238, 3,226]s overlap heavily; we draw no ordering conclusion.

E2b no-mutator (preliminary, $N=3$). With the recipe-guided mutator disabled but the controller, LLM proposal layer, and Ghidra static-analysis channel retained, the 3 `no-mutator` runs again reach the same 269-edge ceiling, with median `last_find` = 13,548s and median plateau = 930s (descriptive -63.3% vs baseline; Mann–Whitney $U=2$, $p=0.14$; $A_{12}=0.87$). Per-run plateaus are [286, 930, 962]s; bootstrap 95% CI is [286, 962]s. Median `execs_per_sec` is 12,712 (slightly below the baseline median 13,049) and median `cycles_done` is 671, much closer to the baseline median (797) than to the 119–140 cycles observed in either mode that engaged the recipe-guided mutator (`rule-only` and `full-agent`).

Ablation ordering contradicts the “mutator-framework drives the gain” hypothesis. The point-estimate ordering of median plateaus is `no-mutator` (930s) < `rule-only` (1,165s) < `full-agent` (1,384s) < `baseline-afl` (2,532s). If the recipe-guided mutator (with or without an LLM-derived recipe) were the sole driver of the gain, removing it in E2b should have moved the median plateau *up* toward baseline, not down past `rule-only`. The most data-consistent account is therefore: (i) the controller’s plateau detector and its corpus-snapshot reorganization are at least partially responsible for the productive-to-plateau acceleration (all three ablation modes share that machinery, all three reduce plateau versus baseline); (ii) engaging the custom mutator slows `cycles_done` by $\approx 5\text{--}6\times$ without yielding a corresponding plateau-duration benefit on cJSON; and (iii) the LLM proposal layer contributes nothing beyond the default rule recipe on this target (RQ3, §6.4). *All three account-changes are descriptive at this N*: the per-run plateau ranges of E2a [674, 1,540]s and E2b [286, 962]s overlap, the Mann–Whitney p -values are well above 0.05 in every pairwise comparison, and the bootstrap CIs for the three ablation modes overlap with `full-agent`’s CI. The preprint therefore reports the ordering as a hypothesis that the *controller orchestration* is the dominant component on saturated cJSON, to be tested at higher N on a non-saturated target.

E2c no-static-analysis (preliminary, $N=3$). With the Ghidra static-analysis channel disabled but the controller, LLM proposal layer, and recipe-guided mutator retained, the 3 `no-static-analysis` runs reach the same 269-edge ceiling, with median `last_find` = 14,250s and median plateau = 234s (descriptive -90.8% vs baseline; Mann–Whitney $U=5$, $p=0.57$; $A_{12}=0.67$). *However the per-run plateau distribution is bimodal* ([187, 234, 4,454]s; bootstrap 95% CI [187, 4,454]s spans the entire interval), so the point estimate is dominated by a single outlier-vs-cluster structure that $N=3$ cannot resolve. Median `execs_per_sec` is 10,259 (lower than other E2 arms, because the empty-blackboard agents take a few extra prompt cycles before falling back to the default recipe) and median `cycles_done` is 96, the lowest of any mode — the combination of recipe-guided mutator slowdown and longer agent-call overhead.

Cross-ablation pattern (and the controller-only hypothesis it suggests). At $N=3$ the four E2 ablation arms order `no-static-analysis` (234s) < `no-mutator` (930s) < `rule-only` (1,165s) < `full-agent` (1,384s), each shorter than the `baseline-afl` plateau (2,532s). All four share the controller’s plateau detector and corpus-snapshot reorganization; the modes that engage the recipe-guided mutator and the LLM proposal layer fall *between* the controller-only-shaped floor (`no-static-analysis` median 234s) and the fully loaded `full-agent` median. The most natural read is that the productive-to-plateau acceleration on cJSON is driven primarily by controller machinery, with the LLM proposal and recipe-guided mutator layered on top adding latency rather than coverage — a hypothesis the venue **controller-only** ablation will test directly. *All of this is descriptive at $N=3$ with overlapping bootstrap CIs*; no pairwise difference between the four ablation arms reaches $p<0.05$ under Mann–Whitney U at this N .

Missing controller-only arm (venue follow-up). The current E2 design does not include a **controller-only** condition (plateau detector + corpus snapshot only, no LLM, no recipe-guided mutator, no Ghidra). The venue version will add it. Without it we cannot fully separate “controller machinery” from “default rule recipe”, and the controller-driven attribution above remains a hypothesis rather than a measured effect. The full multi-arm comparison (`baseline` / `rule-only` / `no-mutator` / `no-static-analysis` / `controller-only` / `full-agent`) at $N\geq 10$ on a non-saturated target is explicit follow-up work (§8).

E4 fairness baselines: AFL++ with dict / cmplog ($N=3$). The preprint `baseline-afl` arm is vanilla AFL++ without `-x dictionary` or `cmplog` instrumentation (§8), so any plateau delta against it overstates FuzzPilot’s advantage. E4 is the corresponding fairness arm: `baseline-afl-dict` runs vanilla AFL++ with `-x experiments/targets/cjson/fuzzpilot_default.dict` (the same FUZZ/MAGIC/TOKEN tokens the controller emits in its default recipe), and `baseline-afl-cmplog` runs vanilla AFL++ over a `cmplog`-instrumented build (`cjson_fuzzer.cmplog`). 3 runs each, 14,400 s, same host. Two findings stand out:

- **Plateau: AFL++ with `-x` or `cmplog` does *not* close the plateau gap.** Medians are `baseline-afl-dict` = 2,256 s ($p=0.79$, $A_{12}=0.60$ vs baseline) and `baseline-afl-cmplog` = 2,367 s ($p=1.00$, $A_{12}=0.53$ vs baseline) — neither is meaningfully different from the 2,532 s vanilla baseline, and both are much longer than the FuzzPilot ablation arms (range 234–1,384 s). The plateau-recovery delta we attribute in §6.5 to the FuzzPilot controller *does not reduce* to “vanilla AFL++ once it has a dictionary” or “vanilla AFL++ once it has runtime I2S information”.
- **Edges: `cmplog` actually reaches one more edge than FuzzPilot does on cJSON.** All three `baseline-afl-cmplog` runs hit 270 edges (vs 269 for `baseline-afl`, `full-agent`, and all E2 ablations; one `baseline-afl-dict` run also hits 270). The +1-edge gap is small but reproducible and *must be disclosed*: `cmplog`’s runtime comparison-operand mining finds an edge that the recipe-guided mutator + default rule recipe does not. The two techniques operate on different axes — FuzzPilot accelerates plateau via controller-driven corpus-snapshot reorganization (saves wall-time), while `cmplog` marginally extends absolute coverage via I2S mining (finds one more edge). They are complementary, not substitutes; composing the two (FuzzPilot’s controller + `cmplog`-instrumented target) is an obvious follow-up that the current FuzzPilot CLI does not yet support.

The plateau-delta narrative we report in §6.2 therefore stands as a wall-time advantage on a target whose absolute coverage is saturated; on a non-saturated target the E4 fairness comparison must be re-run, and we expect the controller-vs-`cmplog` ordering on absolute edges to be the open question.

6.6 Headline summary (consolidated)

For a single-glance reading, Table 11 pulls together every headline RQ number, with the matching significance companion or null-result marker where applicable. Cells marked *desc.* are descriptive point estimates at the listed N ; p is the two-sided Mann–Whitney U test (§6.2) unless noted; CI is the percentile bootstrap 95% over 10,000 resamples. Negative findings (zero promotions, TOST not established, etc.) are explicitly listed rather than buried.

7 Related Work

In-loop LLM fuzzers. LLAMAFUZZ [21] demonstrated that an LLM fine-tuned on seed/mutation pairs can produce structurally valid mutations and consistently outperforms AFL++ on Magma. The model is queried per mutation; throughput is bound by model latency. FuzzCoder [15] casts byte-level mutation as a sequence-to-sequence task. Fuzz4All [24] treats the LLM as the primary input generator and mutation engine across multiple input languages and targets. Semantic-Aware Fuzzing [2] explicitly notes the throughput bottleneck of in-loop designs and offers it as motivation

RQ	Metric	Result (N)	Companion / caveat
RQ1	median plateau, baseline	2,532 s ($N=5$)	CI [539, 3,970]
RQ1	median plateau, full-agent	1,384 s ($N=5$)	CI [238, 3,226]
RQ1	Δ plateau (full vs baseline)	-45.3% (<i>desc.</i>)	MW $U=8$, $p=0.42$
RQ1	median time-to-269 edges, baseline	2,524 s ($N=5$)	range [1,081, 4,448]
RQ1	median time-to-269 edges, full-agent	6,702 s ($N=4/5$)	one run never reached
RQ2	median exec/sec ratio (full / baseline)	$\approx 1.06\times$ ($N=5$)	leave-one-out (dropped)
RQ2	microbench fp-active / fp-empty, median	$1.016\times$ ($N=5$)	mean ratio 0.87x
RQ2	microbench fp-active / fp-empty, t -test	two-sided $p=0.45$	no significant difference
RQ2	TOST equivalence at $\pm 5\%$	$p_{\text{TOST}}=0.68$ ($N=5$)	equivalence not met
RQ2	microbench fp-empty / vanilla, median	$1.53\times$ ($N=5$)	sanity gate ($\leq 5\times$)
RQ3	schema-valid LLM responses	43/45 = 95.6% ($N=5$)	rule-only / no-mutator
RQ3	LLM-recipe promotions (gate accepted)	0/20 ($N=5$)	cJSON saturated, no JSON
RQ3	DeepSeek API cost	\$0.002 / 4h campaign	FX rate \$1 $\approx$ ¥7.2;
RQ4	median plateau, rule-only (E2a)	1,165 s ($N=3$)	MW $U=4$, $p=0.39$
RQ4	median plateau, no-mutator (E2b)	930 s ($N=3$)	MW $U=2$, $p=0.14$
RQ4	median plateau, no-static-analysis (E2c)	234 s ($N=3$)	MW $U=5$, $p=0.57$
RQ4	ablation ordering	no-static-analysis < no-mutator < rule-only < full-agent	all four shorter
RQ4	controller-only ablation	—	not in preprint manuscript
Fair	E4a median plateau, baseline-afl-dict	2,256 s ($N=3$)	$p=0.79$, $A_{12}=0.60$
Fair	E4b median plateau, baseline-afl-cmplog	2,367 s ($N=3$)	$p=1.00$, $A_{12}=0.53$
Fair	E4b edges, baseline-afl-cmplog	270/270/270	cmplog finds +1

Table 11: One-glance summary of all RQ-level numbers reported in the preprint, with the matching significance / null-result companion. *Desc.* = descriptive point estimate; no pairwise plateau or last.find difference reaches $p<0.05$ at this N . Reproducible from `scripts/paper01/review_stats.py`.

for future work. FuzzPilot diverges from this family by treating the LLM as a strategy proposer rather than a per-mutation step.

Off-hot-path placement: G²FUZZ. G²FUZZ [14] is the closest prior work and the one we explicitly position against. Both designs: invoke the LLM only when the fuzzer fails to make progress; have the LLM emit a reusable intermediate rather than per-input mutations; and let AFL++ run at native speed afterwards. The differences are intentional and orthogonal:

- **Target domain.** G²FUZZ targets non-textual binary formats (JPEG, TIFF, MP4, ...); FuzzPilot is *designed for* structured-text parsers and *this preprint evaluates only cJSON*. Generalization claims (XML, SQL, libpng, Magma suite) are deferred to the venue paper.
- **Intermediate representation.** G²FUZZ emits Python generator scripts (code); FuzzPilot emits schema-validated recipes (data). §4.5 compares the two.
- **Validation.** G²FUZZ’s generated inputs enter the AFL++ queue and are filtered by coverage feedback as usual; FuzzPilot’s recipes are pre-promotion-gated by an explicit N -second micro-campaign with a reward function.
- **Static context.** G²FUZZ uses an LLM-extracted “feature list” from format documentation; FuzzPilot uses Ghidra headless on the target binary (function summaries, dictionary candidates, protected ranges), supporting closed-source targets.

We do not attempt a head-to-head numerical comparison against G²FUZZ in this preprint. Their evaluation is on non-textual binary formats outside our scope; ours is on structured-text parsers outside theirs. §8 states this scope split explicitly.

Hybrid static-and-LLM fuzzers. The hybrid fuzzer of Lin et al. [12] integrates static analysis (CFG and data-flow extracted via LLVM IR or angr) with LLM-guided input mutation through a helper fuzzer running at ≈ 250 queries/hour. It also reports a syntax-schema validator and a semantic novelty score. FuzzPilot differs in two ways: (a) the LLM is genuinely off the hot path, not on a parallel hot path, and (b) the validation step is a budgeted AFL++ run with a reward function, not an embedding-distance score against past inputs.

LLM-synthesized mutators for compilers. MetaMut [16] and Mut4All [27] use LLMs to synthesize *static* custom mutators (in C/C++) for compiler fuzzing, compiled once and used many times. FuzzPilot’s recipes are dynamic per-plateau artifacts rather than compiled mutators, and target consumer parsers rather than compilers.

Multi-agent fuzz frameworks. MALF [1] introduces a multi-agent LLM framework for industrial-control-protocol fuzzing with RAG and QLoRA fine-tuning. The role split is similar in spirit to FuzzPilot’s coordinator / plateau-diagnosis / scheduler / cmp-hint / dictionary / format / mutator / corpus / crash-triage agents; the domain (industrial protocols vs structured text parsers) and the intermediate representation (input sequences vs recipes) differ.

LLM-as-generator (whole inputs). TitanFuzz [6] demonstrates that an LLM can synthesize *entire program-level inputs* (Python snippets exercising deep-learning libraries) zero-shot, and that infill-style prompting matches or exceeds template-driven generators on coverage. FuzzGPT [7] extends the idea by prompting the LLM with historical bug-triggering inputs to push it toward edge cases. ChatFuzz [9] couples ChatGPT-style mutations with greybox feedback for structured-text formats (JSON / XML / HTML) and is the closest single-target comparator to FuzzPilot’s structured-text scope. All three generate *inputs* directly; FuzzPilot’s intermediate is a *recipe* that a downstream AFL++ mutator consumes, so the LLM inference cost is amortized across the entire plateau interval rather than per input.

LLM with white-box / static context. WhiteFox [25] drives compiler fuzzing with LLM-synthesized programs that target specific compiler-optimization rules extracted from source by another LLM pass; the white-box context source is the compiler’s own implementation. KernelGPT [26] mines kernel headers and source to build syscall specifications and then has an LLM synthesize Syzlang-like input programs. FuzzPilot’s Ghidra channel (§5.2) is a parallel design point that opts for binary-level static analysis (so the same pipeline applies to closed-source targets) rather than source-level extraction; the trade-off is coarser per-function summaries in exchange for applicability beyond the open-source assumption.

Industrial LLM-driven fuzzing infrastructure. OSS-Fuzz-Gen [13] is Google’s production framework for using LLMs to *generate fuzz targets and harnesses* for the OSS-Fuzz fleet, addressing a complementary problem (harness/target authoring at scale) to FuzzPilot’s (in-campaign recipe proposal). The two could be composed: an OSS-Fuzz-Gen-authored harness fed into a FuzzPilot-orchestrated campaign with a Ghidra-informed recipe layer. We do not evaluate this composition here.

Grammar / AST-aware greybox fuzzers (LLM-free). A second family of structured-text fuzzers does not use LLMs at all and obtains its structure awareness from an explicit grammar or AST. AFLSmart [17] extends AFL with an input-model-aware mutation phase that preserves chunk boundaries; Nautilus [4] drives greybox mutation from a context-free grammar with subtree splicing; Superion [22] pairs AFL with an ANTLR-derived AST-mutation operator. Conceptually FuzzPilot’s recipe vocabulary (§4.2) is byte-level rather than AST-level: a recipe can express “insert dictionary token at offset x ”, not “replace the second `value` child of an `object` node”. We acknowledge this as a design limitation (§8) — on grammar-amenable targets, AFLSmart-style chunk-aware mutation or Nautilus-style subtree splicing remains a strong baseline that we do not attempt to outperform in this preprint. Whether an LLM-guided recipe layer composes *on top of* a grammar-aware mutator (e.g., FuzzPilot recipes that target Nautilus subtrees rather than byte offsets) is an explicit follow-up direction.

AFL++ semantic-ish features (cmplog, RedQueen, laf-intel). A third class of techniques exploits the AFL++ instrumentation itself rather than an external grammar or LLM. The `cmplog` module records concrete comparison operands during fuzzing and feeds them back as dictionary-overwrite candidates; RedQueen [5] pioneered the underlying input-to-state-correspondence (I2S) idea; the `laf-intel` LLVM pass [11] splits multi-byte comparisons into single-byte ones to make magic numbers reachable by coverage feedback alone. These are precisely the AFL++ features a fair **baseline-afl** arm should engage. Our preprint **baseline-afl** arm *does not* pass `-x dictionary` or build the target with `cmplog`, and so the plateau-delta measured against it includes a non-trivial dictionary-contribution component that an AFL++-with-`cmplog` baseline could close. We treat this as the single most important missing baseline (§8); the E4 fairness arms reported in this paper (**baseline-afl-dict** with the FuzzPilot default dictionary, and **baseline-afl-cmplog** with a `cmplog`-instrumented `cjson_fuzzer.cmplog`) close this directly — and find that neither closes the plateau gap, but `cmplog` reaches +1 edge that FuzzPilot does not (see §6.5).

LLM-upfront generators (one-shot, no steady-state LLM calls). Sphinx [19] uses an LLM *once, before fuzzing starts* to synthesize a structured input generator for an SMT solver; the generator then runs without any further LLM involvement. The design rhymes with FuzzPilot in placing the LLM outside the per-mutation hot path, but the LLM is consulted only once (not on every plateau) and the output is a complete generator rather than a recipe that re-tunes a host fuzzer’s existing mutation operators. Sphinx and FuzzPilot together suggest a spectrum — *LLM once* (Sphinx), *LLM at plateaus* (FuzzPilot, G²FUZZ), *LLM per input* (LLAMAFUZZ, ChatFuzz) — that we expect to be a productive axis for future comparisons.

T3: side-by-side comparison.

8 Limitations, Threats to Validity, and Future Work

Scope: cJSON only in this preprint. This preprint evaluates FuzzPilot on cJSON only (5 **baseline-afl** + 5 **full-agent** runs at 14,400s, plus the E3 microbench and 3 preliminary **rule-only** runs from E2a). The architecture is designed for structured-text parsers, but no empirical claim is made about libpng, libxml2, sqlite3, openssl, or any non-textual binary format here; on those domains G²FUZZ’s code-as-generator is plausibly stronger. A head-to-head comparison on G²FUZZ’s benchmark (JPEG, TIFF, MP4, etc.) and the larger structured-text targets are deferred to the venue version of the paper.

cJSON is a saturated target for the LLM proposal layer. Within the 14,400 s budget the baseline AFL++ harness reaches the 269-edge ceiling at a per-run median of $\approx 2,500$ s (slowest baseline run 4,448 s), leaving the parent corpus at plateau time with no edge headroom for a 20 s micro-campaign to find. As a result, the LLM proposal layer’s contribution (the *LLM-promoted recipe* path, distinct from the *mutator framework*) cannot be empirically demonstrated on this target: every micro-campaign reward is $R=0$ and every promotion is skipped (§6.4). A non-saturated target where the gate has edge headroom is required to evaluate the LLM layer in isolation; this is the most important single piece of future work for the venue paper.

Ghidra benefit not fully demonstrated. Our targets are open-source; LLVM-IR or angr-based static analysis would, in principle, supply the same blackboard fields used here. The motivation for Ghidra is to make the architecture applicable to closed-source binary targets, and the E2c ablation (§6.5) only quantifies the contribution of static context, not the specific advantage of Ghidra over alternatives. A binary-target evaluation that isolates Ghidra’s contribution is explicit future work.

Repetition count and statistical power. We run 3–5 repeats per cell, which is consistent with prior LLM-fuzzing preprints [12, 14, 21] but insufficient for formal significance testing. We report medians, interquartile ranges, and per-run raw numbers rather than p -values. The venue paper will use ≥ 10 repeats with rank-sum tests.

Crash deduplication not in scope. We report coverage and plateau-recovery, not unique-bug counts. Crash deduplication (e.g. the casr-cluster pipeline) is deferred.

Deferred ablations and sensitivity. The *ai-direct* ablation (promote every schema-valid proposal without micro-campaign), the *random-recipe* ablation (skip the LLM entirely and emit randomly-sampled recipes), and a micro-campaign-parameter sensitivity sweep (over N , K , and the reward weights $\alpha, \beta, \gamma, \delta_h, \delta_m$) are wired into the CLI but not in the manifest for the preprint matrix. Their data lands in the venue paper, alongside the remaining *no-static-analysis* arm of the E2 ablation (the *rule-only* and *no-mutator* arms are already reported in §6.5).

Recipe vocabulary is byte-level, not AST-level. The seven-operator recipe vocabulary (*BitFlip*, *OverwriteRange*, *InsertToken*, *Arith*, *Splice*, *DeleteBlock*, *DictionaryOverwrite*; §4.2) operates on *byte offsets within a flat input buffer*, not on syntactic nodes of the structured-text grammar. A recipe can say “insert the token “null” at offset 128”; it cannot say “replace the second child of an *object* node with null”. For grammar-amenable targets, AST-level operators (*Nautilus* [4], *Superion* [22]) or chunk-level mutations (*AFLSmart* [17]) are stronger generators of parser-valid mutants. We chose a byte-level vocabulary deliberately — to keep the mutator hot-path simple and to keep the recipe schema small enough for the LLM to emit reliably — but the trade-off is real, and *FuzzPilot* is unlikely to *outperform* a grammar-aware fuzzer on grammar-amenable targets without additional integration. Composing *FuzzPilot*’s recipe layer on top of a grammar-aware mutator (recipes that target AST subtrees rather than byte offsets) is an explicit future-work direction.

No syntax-aware repair or invariant-preserving operators. Several recipe operators (*OverwriteRange*, *DictionaryOverwrite*, *Splice*) can produce ill-formed JSON if applied naively: an overwrite that straddles a string-literal boundary breaks the quote-pair invariant; a splice that lands inside a numeric literal can produce an unparseable token; a dictionary overwrite inside a

key context can violate the keys-are-unique invariant. The current recipe schema has *no length-preserving, quote-aware, or escape-aware* repair operator. AFL’s downstream coverage feedback eventually discards inputs that fail to advance coverage, so the practical impact on the campaign is bounded; but a sharper grammar-aware mutant generation strategy is the obvious next step. We have considered adding a `LengthRepair` operator (rebalance quote/bracket pairs after a mutation) and a `StringInsert` operator (insert at safe byte boundaries within a string literal), but neither is implemented in the released version — avoiding having to specify them is part of why we kept the schema small and audit-friendly. We acknowledge this is a real reviewer-grade limitation rather than a design victory.

AFL++ fairness baseline (-x dictionary and cmplog) is a follow-up arm, not finished in this preprint. The `baseline-afl` arm reported in §6.2 is *vanilla* AFL++: it does *not* receive an `-x dictionary` file, and the target binary is *not* compiled with the `cmplog` pass. FuzzPilot’s recipe-guided mutator, by contrast, dispatches against the controller’s default `DictionaryAgent` recipe whose token set is `{"FUZZ", "MAGIC", "TOKEN"}` (§6.4). This asymmetry overstates FuzzPilot’s advantage: any plateau-recovery delta could in principle be closed by an AFL++ run that consumes the *same* three tokens via `-x`, or by an AFL++ build that uses `cmplog`/RedQueen-style I2S mining [5, 11] to infer target-specific tokens at runtime. We treat closing this asymmetry as the single most important missing arm in the preprint matrix. The E4 fairness baselines (`baseline-afl-dict` with the FuzzPilot default dictionary; `baseline-afl-cmplog` with a `cmplog`-instrumented `cjson_fuzzer.cmplog`) are reported in §6.5; the headline finding is that neither closes the plateau gap, but `cmplog` reaches one more edge (270) than any FuzzPilot arm (269). §9 pins the dict file and the `cmplog` binary so reviewers can re-run the comparison themselves.

Reward / hyperparameter sensitivity is fixed, not swept. The micro-campaign budget $N=20$ s, the candidate count $K_{\text{cand}}=4$, and the reward weights $(\alpha, \beta, \gamma, \delta_h, \delta_m)$ in §5.3 are pinned to fixed values for the preprint matrix — chosen during development on the cJSON corpus, then frozen. We did not run a sensitivity sweep. The risk surface that this exposes is twofold. (i) $N=20$ s may be too short for the gate to discriminate marginal candidates on a non-saturated target (the cJSON null result is agnostic on this since reward is identically zero). (ii) The reward weights bias the gate toward edge-discovery; on a crash-finding-dominant target, a different $(\alpha, \delta_h, \delta_m)$ trade could matter. A small sensitivity sweep (e.g., $N \in \{10, 20, 40, 80\}$ s and $K_{\text{cand}} \in \{2, 4, 8\}$ at $N=3$ each) is wired into the CLI and is part of the venue-paper matrix.

Multi-target generalization (venue rollout plan). We do not run on Magma’s full benchmark suite in this preprint, and we do not claim generalization across target families. The venue paper adds three deeper structured-text / binary-schema targets that exercise components of FuzzPilot that cJSON’s 269-edge ceiling does *not*: `libxml2` (XML parser, ~30k edges, AFLS-mart/Nautilus/Superion comparator), `sqlite3` (SQL parser/VM, ~150k edges, OSS-Fuzz flagship), and `openssl_x509` (ASN.1 + X.509 parser, ~80k edges, binary + deep-schema). The venue evaluation matrix (`experiments/manifests/paper01_venue.yaml`) follows Klees et al. [10] at $N \geq 20$ main runs and $N \geq 10$ ablation runs per target, 24h budget per run (Magma/OSS-Fuzz convention). A Stage-1 pilot ($N=5$ main / $N=3$ ablations, 24h budget) is currently in flight on the experiment host with the harnesses (`libxml2_fuzzer`, `sqlite3_fuzzer`, `x509_fuzzer`) and their per-target Ghidra dictionaries checked into the repository; the Stage-1 outcome gates whether the Stage-2 full-Klees matrix runs on all three targets or focuses on the subset where the LLM proposal layer shows a non-null contribution. *Headroom check*: a 60s smoke-test run of vanilla `afl-fuzz` on the

three harnesses (no FuzzPilot pipeline) reaches 2,776 edges on libxml2, 6,922 edges on sqlite3, and 3,635 edges on openssl X.509 — 10–25 $\times$ more than cJSON’s 269-edge ceiling at the *end* of a 14,400 s budget. The Stage-2 budget will not saturate any of the three, so the LLM proposal layer has room to contribute a non-null result if it can. The choice to ship the preprint at cJSON-only is deliberate: it lets reviewers verify the infrastructure (audit, gate, mutator, micro-campaign) end-to-end on a target small enough to read every artifact, before we burn $\sim$ 1,000 CPU-hours per target on the Stage-2 matrix. *No preprint claim depends on the venue targets*; their data lands in v2.

Platform. The development platform is macOS/arm64 (with a Docker linux/arm64 image); the canonical evaluation platform is Linux/x86_64 (a 4-core Intel Xeon VM with a Linux/amd64 native build). All numbers in this paper are from the canonical platform.

LLM nondeterminism. Even with temperature pinned at 0.0, model providers do not guarantee bit-exact reproducibility across versions. We mitigate by (a) recording the audit trail so any decision is replayable from the recipe rather than the model, (b) reporting schema-validity and fallback rates per run, and (c) releasing a pinned Docker image that captures the FuzzPilot tag and the model client version at submission time.

Threats to validity. We organize the threats following Wohlin et al.’s [23] standard internal / external / construct / conclusion taxonomy, with an additional explicit discussion of LLM-specific risks.

Internal validity. The plateau-detector hyperparameters ($W=10$ s, $\theta_e=50$, $\theta_p=1$) and the micro-campaign reward weights ($\alpha, \beta, \gamma, \delta_h, \delta_m$) were tuned on a *development cJSON seed corpus* (the 26 hand-curated JSON files described in §6.1) that became the *evaluation* seed corpus once the design was frozen. We did not implement a hold-out split between tuning and evaluation corpora and therefore cannot exclude data leakage; the venue version will tune on a disjoint development corpus or via cross-validation across seed-corpus splits, and will report the delta from tuning-set to evaluation-set numbers. A second internal threat is the absence of a **controller-only** ablation (§6.5) that would separate the plateau detector / corpus-snapshot reorganization machinery from the default rule recipe; we note this as a known design gap to be addressed in the venue ablation matrix. Plateau-detector SIGSTOP-based microbench isolation (§5.3) is also a methodological hazard: running the microbench concurrently with the main AFL workers *without* SIGSTOP inverts our central RQ2 throughput-parity ratio in pilot tests, so any reproducer that omits the isolation will obtain different numbers.

External validity. A single target family (cJSON, parser of a structured-text format) is not representative; we make *no* cross-format generalization claim and explicitly defer it. The LLM proposal layer in particular is exercised but not stress-tested by the cJSON target: the gate returns 20/20 null candidates (§6.4), so we observe only the *no-false-positive* half of the gate’s intended behavior. A single LLM (**deepseek-chat**) at temperature 0 is also a single-point evaluation in model space; whether **claude-haiku-4-5**, **gpt-4o-mini**, **llama-4-maverick**, or any other backend produces a materially different schema-validity rate, recipe distribution, or null-result behavior is open. Finally, all numbers were collected on a single 4-vCPU experiment host (single tenancy at collection time except where noted in §6.3); cross-machine reproducibility is supported by the Docker recipe but not separately measured.

Construct validity. The choice of metric is itself a construct decision. We use plateau duration (**run_time** – **last_find**) as a primary endpoint metric, which is sensitive to a single late-budget discovery (e.g. full-agent r04 finds the 269th edge at 14,205 s and is credited with a 238 s plateau; see

Table 5). Klees et al. [10] recommend coverage-AUC or time-to- N -edges as endpoint-free alternatives; we therefore additionally report time-to- N -edges for $N \in \{264, 268, 269\}$ in Table 4, which gives the opposite headline (**baseline-afl** reaches the 269-edge ceiling *faster* than **full-agent** at the median). The two metrics are mutually consistent — a plateau-shortening intervention can also delay the ceiling-touch by spreading out the discovery curve — but the construct choice matters and we report both. The 269-edge “ceiling” is itself a construct: it is the number of AFL++ bitmap slots hit by the cJSON harness, not a measure of all reachable basic-block edges in the cJSON library; a richer harness (or a different fuzzing instrumentation) could expose more edges.

Conclusion validity. With $N=5$ per main arm and $N=3$ per ablation, we are well below the $N \geq 20$ guideline of Klees et al. [10]. We therefore report all group differences as descriptive point estimates accompanied by Mann–Whitney U , exact p -values, Vargha–Delaney A_{12} [20], and percentile-bootstrap 95% confidence intervals (see §6.2). No descriptive delta in this preprint reaches $p < 0.05$ on the two-sided Mann–Whitney test; we present each result as a hypothesis rather than as a tested claim and the venue version will repeat the experiments at $N \geq 10$ with formal significance testing.

LLM-specific reproducibility. The model identifier **deepseek-chat** is a moving alias served by DeepSeek; the underlying weights may change without notice between runs. At $T=0$ the API surface is intended to be deterministic, but our data shows 2/45 schema-invalid responses, indicating that the API surface is not bit-exact deterministic in practice. We mitigate by (a) logging the **model** field returned in each API response into **agent_decisions.jsonl**, (b) recording the full prompt / response pair so any decision is replayable from the recipe rather than the model, and (c) shipping a Docker recipe that pins the SDK version (not the model version). A future-proof reproducibility strategy would also archive the exact prompt-response transcripts in a content-addressable store to allow regeneration without re-querying the API.

9 Conclusion

FuzzPilot adapts the off-hot-path LLM-control paradigm pioneered by G²FUZZ from non-textual binary formats to the structured-text-parser setting by building three reusable pieces of infrastructure: a data-as-recipe intermediate that replaces LLM-generated code, a micro-campaign validation gate that scores candidate recipes against an explicit reward function before promotion, and a Ghidra-derived static context channel that does not enter the mutation hot path. On cJSON — the single target evaluated in this preprint — the **baseline-afl** and **full-agent** groups both reach the same 269-edge target ceiling within their 14,400s budgets, the median **execs.per.sec** ratio of **full-agent** to baseline is $\approx 1.06\times$, and per-run median plateau duration is a descriptive 45.3% shorter under **full-agent** (1,384s vs 2,532s; Mann–Whitney $U=8$, $p=0.42$ at $N=5$ — significance is therefore *not* established at this N). Across the 5 **full-agent** runs the micro-campaign gate evaluated 20 LLM-derived candidates and *promoted zero*, because cJSON’s 269-edge ceiling is reached well before plateau time and the gate’s reward is identically $R=0$. The plateau-duration delta is therefore attributable to the recipe-guided mutator running the controller’s default **DictionaryAgent** recipe rather than to LLM-promoted recipes; the ablation ordering in §6.5 (**no-mutator** < **rule-only** < **full-agent** on plateau median) further suggests the controller’s plateau detector and corpus-snapshot reorganization are co-drivers, though this remains a descriptive hypothesis at $N=3$ ablations and is not a measured effect. The LLM proposal layer’s positive coverage contribution cannot be empirically demonstrated on cJSON; what *is* demonstrated is that the infrastructure behaves predictably (43/45 = 95.6% schema-valid, 0 fallback, \$0.002 per 4-hour campaign), so the audit and cost properties of the design hold even though the proposal layer’s coverage benefit does not. An

evaluation on non-saturated targets where the gate has edge headroom is explicit future work in the venue paper. The full implementation, configurations, decision logs, and a reproducibility appendix recording the exact commit, AFL++ version, model configuration, and per-run artifact paths are released alongside the paper. Larger-target evaluation (libpng, libxml2, sqlite3, openssl, Magma) and the remaining `no-static-analysis` ablation row are explicit future work; the `rule-only` (E2a) and `no-mutator` (E2b) preliminary numbers are summarized in §6.5.

Reproducibility appendix

The numbers in this preprint are reproducible from the following pinned artifacts. All paths are relative to the repository root.

- **Repository commit:** 85223d8 (branch `main`, repo https://github.com/Qiao-Zhiyi/fuzz_agent). The arXiv-v1 commit pin is the literal SHA above; the `paper01-arxiv-v1` git tag (created at submission time) points at that same commit so reviewers can fetch it without scrolling history.
- **Fuzzer:** AFL++ build identifier `++4.21c` (the leading `++` is the upstream banner prefix; the underlying release tag is `v4.21c`). Built from source on the experiment host and confirmed in every `fuzzer_stats` file under `results/paper01_ai_recipe_mutation/runs/`.
- **Target binary:** cJSON parser harness at `experiments/targets/cjson/`, AFL persistent mode (`persistent shmем_testcase deferred`). The target binary SHA-256 and the seed-corpus tarball SHA-256 are written per-run into `run_metadata.json` (fields `target_sha256` and `seed_corpus_sha256`); reviewers can verify reproduction by comparing those hashes across re-builds.
- **Seed corpus:** 26 hand-curated JSON files at `experiments/targets/cjson/seeds/`; the directory is small enough to inspect directly and is included in the repository.
- **LLM:** deepseek-chat via the `api.deepseek.com` OpenAI-compatible endpoint (configuration in `experiments/targets/cjson/config.yaml`), temperature 0.0. *Note that `deepseek-chat` is a moving alias served by DeepSeek; the specific underlying model version is logged into the `model` field of each `agent_decisions.jsonl` record at call time, so a reviewer can recover the model-version string from the audit log.* A NIM-hosted `meta/llama-4-maverick-17b-128e-instruct` fallback path is wired through the same watchdog and was exercised in regression testing only — all numbers in this paper are from `deepseek-chat`.
- **Build mode:** native x86_64 Linux on the experiment host (no container in the hot path). The repository ships a Docker recipe at `docker/ubuntu/Dockerfile` (build arg `AFLPP_REF=v4.21c`) for portability. The image digest at submission time is recorded in `docker/ubuntu/IMAGE_DIGEST.txt` once the Docker image is published; the native-build numbers in this paper are the reference. We do *not* claim bit-identical reproduction between the Docker and native builds; the Docker recipe is shipped to lower the activation energy for re-running the pipeline rather than to certify exact reproduction.
- **Per-run artifacts:** 25 completed 14,400s runs (5 `baseline-afl`, 5 `full-agent`, 3 `rule-only`, 3 `no-mutator`, 3 `no-static-analysis`, 3 `baseline-afl-dict`, 3 `baseline-afl-cmplog`) are checked into the repository under `results/paper01_ai_recipe_mutation/runs/`; this is the complete preprint matrix on cJSON. Each run directory contains the unmodified

AFL++ fuzzer_stats, coverage.csv, events.jsonl, agent_decisions.jsonl (full-agent / no-mutator modes), main_launch.sh (the exact command-line that started the AFL run), run_metadata.json, and git.patch (the working-tree diff at run start) for byte-level reproducibility of the experiment.

- **Per-mode launch commands.** Each ablation mode is selected by a single `--ablation` switch to the controller; the canonical invocations are:

```
# E1a baseline-afl
python3 -m fuzzpilot.runner --target cJSON \
--ablation baseline-afl --budget 14400 --run-id p1.e1.cjson.baseline-afl.r01

# E1b full-agent
python3 -m fuzzpilot.runner --target cJSON \
--ablation full-agent --budget 14400 --run-id p1.e1.cjson.full-agent.r01

# E2a rule-only
python3 -m fuzzpilot.runner --target cJSON \
--ablation rule-only --budget 14400 --run-id p1.e2.cjson.rule-only.r01

# E2b no-mutator
python3 -m fuzzpilot.runner --target cJSON \
--ablation no-mutator --budget 14400 --run-id p1.e2.cjson.no-mutator.r01

# E2c no-static-analysis
python3 -m fuzzpilot.runner --target cJSON \
--ablation no-static-analysis --budget 14400 --run-id p1.e2.cjson.no-static-analysis.r01
```

The host-wide orchestrator `scripts/paper01/run_all_host.sh --parallel 4` iterates the full 19-run preprint matrix in 4-way parallel with the same per-mode invocations under the hood.

- **Aggregation:** run `python3 scripts/prepare_paper01_data.py --run-root results/paper01_ai_recipe --out-dir results/paper01_ai_recipe_mutation --manifest experiments/manifests/paper01_preprint` to regenerate `validity_report.md`, `tables/run_level.csv`, and the per-RQ summary tables cited in §6. For the statistical numbers (Mann–Whitney U , A_{12} , bootstrap CIs, time-to- N -edges, and the leave-one-out throughput-parity check), run `python3 scripts/paper01/review_stats.py`; output is deterministic.

Pending for the venue version: **no-static-analysis** (E2c) ablation completion, the missing **controller-only** ablation (§6.5), $N \geq 10$ repeats with formal Mann–Whitney U significance testing per Klees et al. [10], libpng/libxml2/sqlite3 additional targets, and head-to-head measurement against G²FUZZ on its binary-format benchmark. The preprint intentionally limits empirical claims to what the 16 cJSON runs (plus E3 microbench) support.

References

- [1] Anonymous. MALF: A multi-agent LLM framework for intelligent fuzzing of industrial control protocols. *arXiv preprint arXiv:2510.02694*, 2025.
- [2] Anonymous. Semantic-aware fuzzing: An empirical framework for llm-guided, reasoning-driven input mutation. *arXiv preprint arXiv:2509.19533*, 2025.
- [3] Andrea Arcuri and Lionel Briand. A practical guide for using statistical tests to assess randomized algorithms in software engineering. In *Proceedings of the 33rd International Conference on Software Engineering (ICSE '11)*, pages 1–10, 2011.

- [4] Cornelius Aschermann, Tommaso Frassetto, Thorsten Holz, Patrick Jauernig, Ahmad-Reza Sadeghi, and Daniel Teuchert. NAUTILUS: Fishing for deep bugs with grammars. In *Proceedings of the Network and Distributed System Security Symposium (NDSS '19)*, 2019.
- [5] Cornelius Aschermann, Sergej Schumilo, Tim Blazytko, Robert Gawlik, and Thorsten Holz. REDQUEEN: Fuzzing with input-to-state correspondence. In *Proceedings of the Network and Distributed System Security Symposium (NDSS '19)*, 2019.
- [6] Yinlin Deng, Chunqiu Steven Xia, Haoran Peng, Chenyuan Yang, and Lingming Zhang. Large language models are zero-shot fuzzers: Fuzzing deep-learning libraries via large language models. In *Proceedings of the 32nd ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA '23)*, pages 423–435, 2023.
- [7] Yinlin Deng, Chunqiu Steven Xia, Chenyuan Yang, Shizhuo Dylan Zhang, Shujing Yang, and Lingming Zhang. Large language models are edge-case generators: Crafting unusual programs for fuzzing deep learning libraries. In *Proceedings of the 46th IEEE/ACM International Conference on Software Engineering (ICSE '24)*, 2024.
- [8] Andrea Fioraldi, Dominik Maier, Heiko Eißfeldt, and Marc Heuse. AFL++: Combining incremental steps of fuzzing research. USENIX Workshop on Offensive Technologies (WOOT), 2020. Software v4.21c used in this paper.
- [9] Jie Hu, Qian Zhang, and Heng Yin. ChatFuzz: Augmenting greybox fuzzing with large language models. *arXiv preprint arXiv:2308.11525*, 2023.
- [10] George Klees, Andrew Ruef, Benji Cooper, Shiyi Wei, and Michael Hicks. Evaluating fuzz testing. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security (CCS '18)*, pages 2123–2138, 2018.
- [11] laf intel. Circumventing fuzzing roadblocks with compiler transformations. <https://lafintel.wordpress.com/2016/08/15/circumventing-fuzzing-roadblocks-with-compiler-transformations/>, 2016. The laf-intel LLVM pass; integrated into AFL++ as AFL_LLVM_LAF_* build flags.
- [12] Shiyin Lin. Hybrid fuzzing with LLM-guided input mutation and semantic feedback. *arXiv preprint arXiv:2511.03995*, 2025.
- [13] Dongge Liu, Jonathan Metzman, Oliver Chang, and Google OSS-Fuzz team. OSS-Fuzz-Gen: An open framework for LLM-driven fuzz target generation. *arXiv preprint arXiv:2404.14924*, 2024.
- [14] K. Liu et al. Low-cost and comprehensive non-textual input fuzzing with LLM-synthesized input generators. In *Proceedings of the 34th USENIX Security Symposium (USENIX Security '25)*, 2025. Also arXiv:2501.19282. Code at <https://github.com/G2FUZZ/G2FUZZ>.
- [15] Liqun Liu et al. FuzzCoder: Byte-level fuzzing test via large language model. *arXiv preprint arXiv:2409.01944*, 2024.
- [16] Conghua Ou et al. The Mutators reloaded: Fuzzing compilers with large language model generated mutators. In *Proceedings of the 29th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS '24)*, 2024.

- [17] Van-Thuan Pham, Marcel Böhme, Andrew E. Santosa, Alexandru Răzvan Caciulescu, and Abhik Roychoudhury. Smart greybox fuzzing. In *IEEE Transactions on Software Engineering (TSE)*, 2021. Originally NDSS 2019 (AFLSMART); journal extension 2021.
- [18] Donald J. Schuirmann. A comparison of the two one-sided tests procedure and the power approach for assessing the equivalence of average bioavailability. *Journal of Pharmacokinetics and Biopharmaceutics*, 15(6):657–680, 1987.
- [19] Yuyan Sun et al. Sphinx: A language model guided generator for solver fuzzing. In *Proceedings of the International Conference on Automated Software Engineering (ASE '24)*, 2024. Representative “LLM upfront, no steady-state calls” design for SMT-solver fuzzing.
- [20] András Vargha and Harold D. Delaney. A critique and improvement of the CL common language effect size statistics of McGraw and Wong. *Journal of Educational and Behavioral Statistics*, 25(2):101–132, 2000.
- [21] Hongxiang Wang, Xiangwei Xu, Xiaofei Xie, et al. LLAMAFUZZ: Large language model enhanced greybox fuzzing. *arXiv preprint arXiv:2406.07714*, 2024. Latest revision March 2026.
- [22] Junjie Wang, Bihuan Chen, Lei Wei, and Yang Liu. Superion: Grammar-aware greybox fuzzing. In *Proceedings of the 41st International Conference on Software Engineering (ICSE '19)*, 2019.
- [23] Claes Wohlin, Per Runeson, Martin Höst, Magnus C. Ohlsson, Björn Regnell, and Anders Wesslén. *Experimentation in Software Engineering*. Springer Berlin Heidelberg, 2012.
- [24] Chunqiu Steven Xia, Matteo Paltenghi, Jia Le Tian, Michael Pradel, and Lingming Zhang. Fuzz4All: Universal fuzzing with large language models. *arXiv preprint arXiv:2308.04748*, 2024. Originally posted 2023; updated 2024.
- [25] Chenyuan Yang, Yinlin Deng, Runyu Lu, Jiayi Yao, Jiawei Liu, Reyhaneh Jabbarvand, and Lingming Zhang. WhiteFox: White-box compiler fuzzing empowered by large language models. In *Proceedings of the ACM on Programming Languages (OOPSLA '24)*, 2024.
- [26] Chenyuan Yang, Zhou Ruixing Zhao, and Lingming Zhang. KernelGPT: Enhanced kernel fuzzing via large language models. In *Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS '25)*, 2025.
- [27] J. Ye et al. Mut4All: Fuzzing compilers via LLM-synthesized mutators learned from bug reports. *arXiv preprint arXiv:2507.19275*, 2025.

System	LLM placement	LLM output	Validation	Static ctx.	Audit	Target domain
AFL++ [8]	<i>none</i>	—	coverage	none	—	general
AFL++ cmplog / RedQueen [5]	none (I2S)	dict override	coverage	runtime cmp ops	—	general
laf-intel [11]	none (LLVM pass)	—	coverage	cmp splitting	—	general
AFLSmart [17]	none (input model)	chunk mutation	coverage	input-format model	no	structured (PNG / WAV / PDF)
Nautilus [4]	none (grammar)	subtree splicing	coverage	context-free grammar	no	grammar-amenable (JS / SQL / PHP)
Superion [22]	none (grammar)	AST mutation	coverage	ANTLR grammar	no	structured text (XML / JS)
Sphinx [19]	LLM upfront, one-shot	input generator (code)	coverage	none	no	SMT solvers
LLAMAFUZZ [21]	in-loop	input bytes	coverage	none	no	general (Magma suite)
FuzzCoder [15]	in-loop	byte mutation	coverage	none	no	general
Fuzz4All [24]	in-loop	input bytes	coverage	none	no	multi-language
SemAware [2]	parallel helper	input bytes	semantic novelty	none	no	general
Hybrid [12]	parallel helper, ≈ 250 rph	input bytes	novelty + syntax	LLVM-IR / angr	no	libpng / tcpdump / sqlite
TitanFuzz [6]	generator (zero-shot)	whole program input	coverage	none	no	DL libraries (PyTorch / TF)
FuzzGPT [7]	generator (history-prompted)	whole program input	coverage	none	no	DL libraries
ChatFuzz [9]	in-loop (ChatGPT mutator)	input bytes	coverage	none	no	structured text (JSON / XML / HTML)
WhiteFox [25]	two-stage generator	whole program input	coverage	source-level (compiler rules)	no	compilers (PyTorch / LLVM)
KernelGPT [26]	generator + spec-mining	syscall sequences	coverage	source-level (kernel)	no	kernels (Linux)
OSS-Fuzz-Gen [13]	offline (target authoring)	fuzz harness code	build + coverage	source-level (OSS repos)	no	OSS-Fuzz fleet
MetaMut [16]	offline	C/C++ mutator	manual	none	no	compilers
Mut4All [27]	offline	mutator code	manual	bug reports	no	compilers
MALF [1]	multi-agent in-loop	input seqs	coverage	protocol KB	no	industrial protocols
G ² FUZZ [14]	off-hot-path	Python generator (code)	downstream coverage	LLM-extracted features	no	non-textual binary (JPEG / TIFF / MP4)
FuzzPilot (ours)	off-hot-path	schema-validated recipe (data)	<i>N</i>-sec micro-campaign	Ghidra headless	yes	cJSON (preliminary; designed for structured-text parsers)

Table 12: Comparison across LLM-augmented and LLM-adjacent fuzzers. FuzzPilot is positioned as a sibling of G²FUZZ in the off-hot-path quadrant rather than a head-on competitor; the row-by-row differences in intermediate (data vs. code), validation (micro-campaign vs. downstream coverage), context source (Ghidra vs. LLM-extracted features), and target domain (structured text vs. non-textual binary) establish the contribution.